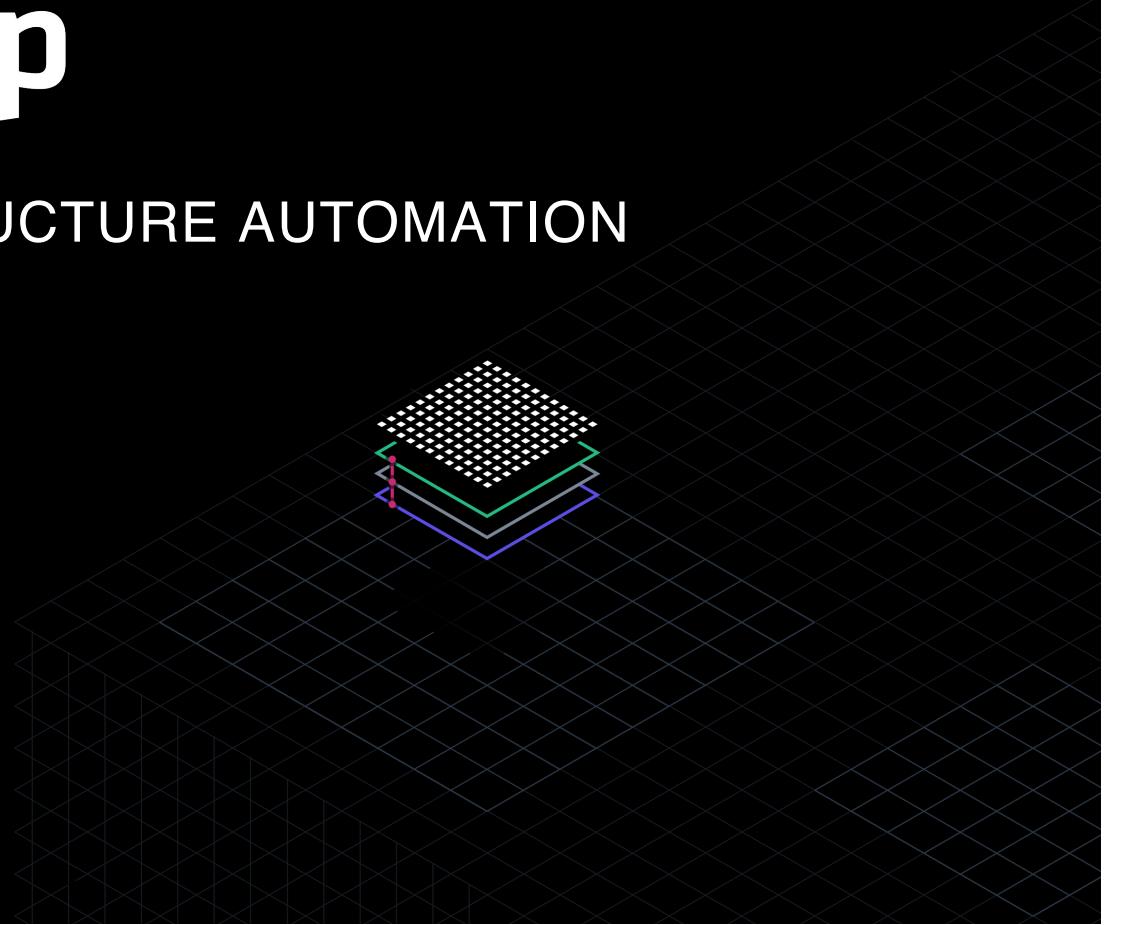
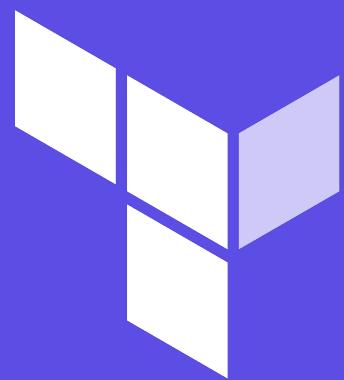




CLOUD INFRASTRUCTURE AUTOMATION





HashiCorp
Terraform

Introduction

Gabe Maentz

- Married with 4 kids (Marion, Victoria, Ellie, Rhys & Ian and new puppy - Zoey)
- Live in the Pittsburgh, PA
- Grew up in Massachusetts, Georgia, Florida, Minnesota and Wisconsin
- Traveling and exploring new cities, soccer, cooking.
- Been working in IT field for almost 20 years: enterprises and startups. (Epic, Kimberly-Clark, Dick's Sporting Goods, DataGravity)
- Run a local community called vBrisket: Fusion between Technology, BBQ and Craft Beer.
- River Point Technology - Specialized HashiCorp Partner
- 7x VMware vExpert, HashiCorp Terraform & Vault Certified



River Point Technology

Tell us about you...

- . Name/Role
- . Company
- . Terraform/Infrastructure as Code experience
- . Which clouds are you working with?
- . Expectations for class

Course Agenda

Terraform Fundamentals

Students will walk away with a solid understanding of HashiCorp Terraform.

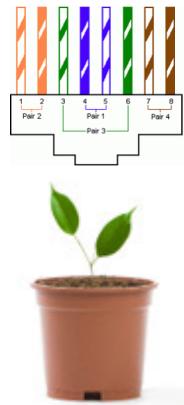
Course Agenda:

- . Evolution of Infrastructure - VCS & Remote Driven
- . Terraform Concepts – Deploy, Init, Apply
- . Terraform Configuration
- . Outputs
- . Console
- . Variables
- . Auto-Formatting
- . Modules
- . Provisioners
- . Meta Arguments
- . Destroy

Progression of Infrastructure Management



Maturity



Manual Everything

```
#!/bin/bash
```



Basic Automation



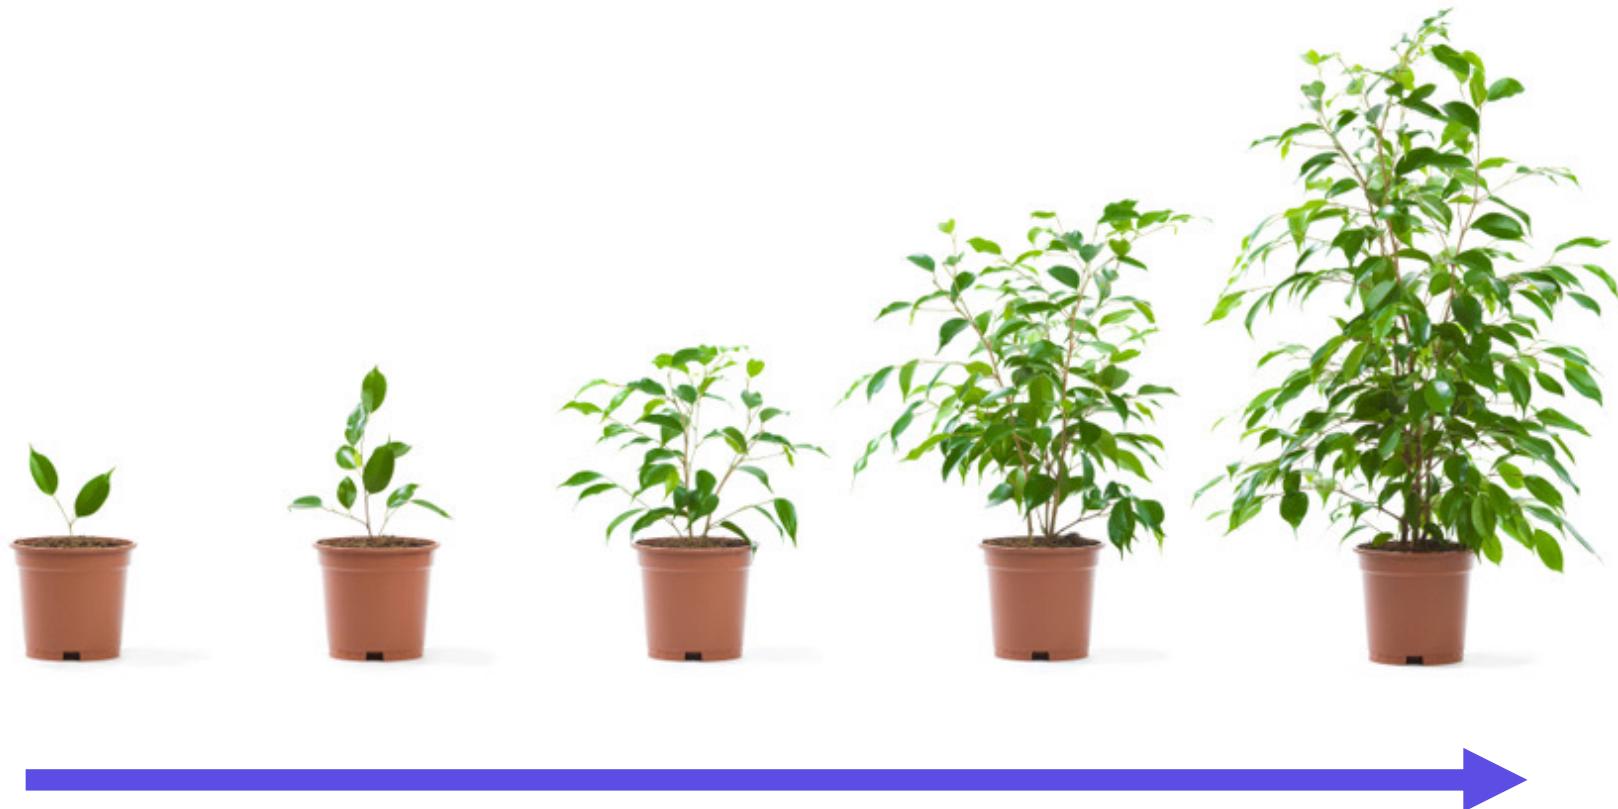
Machine Virtualization



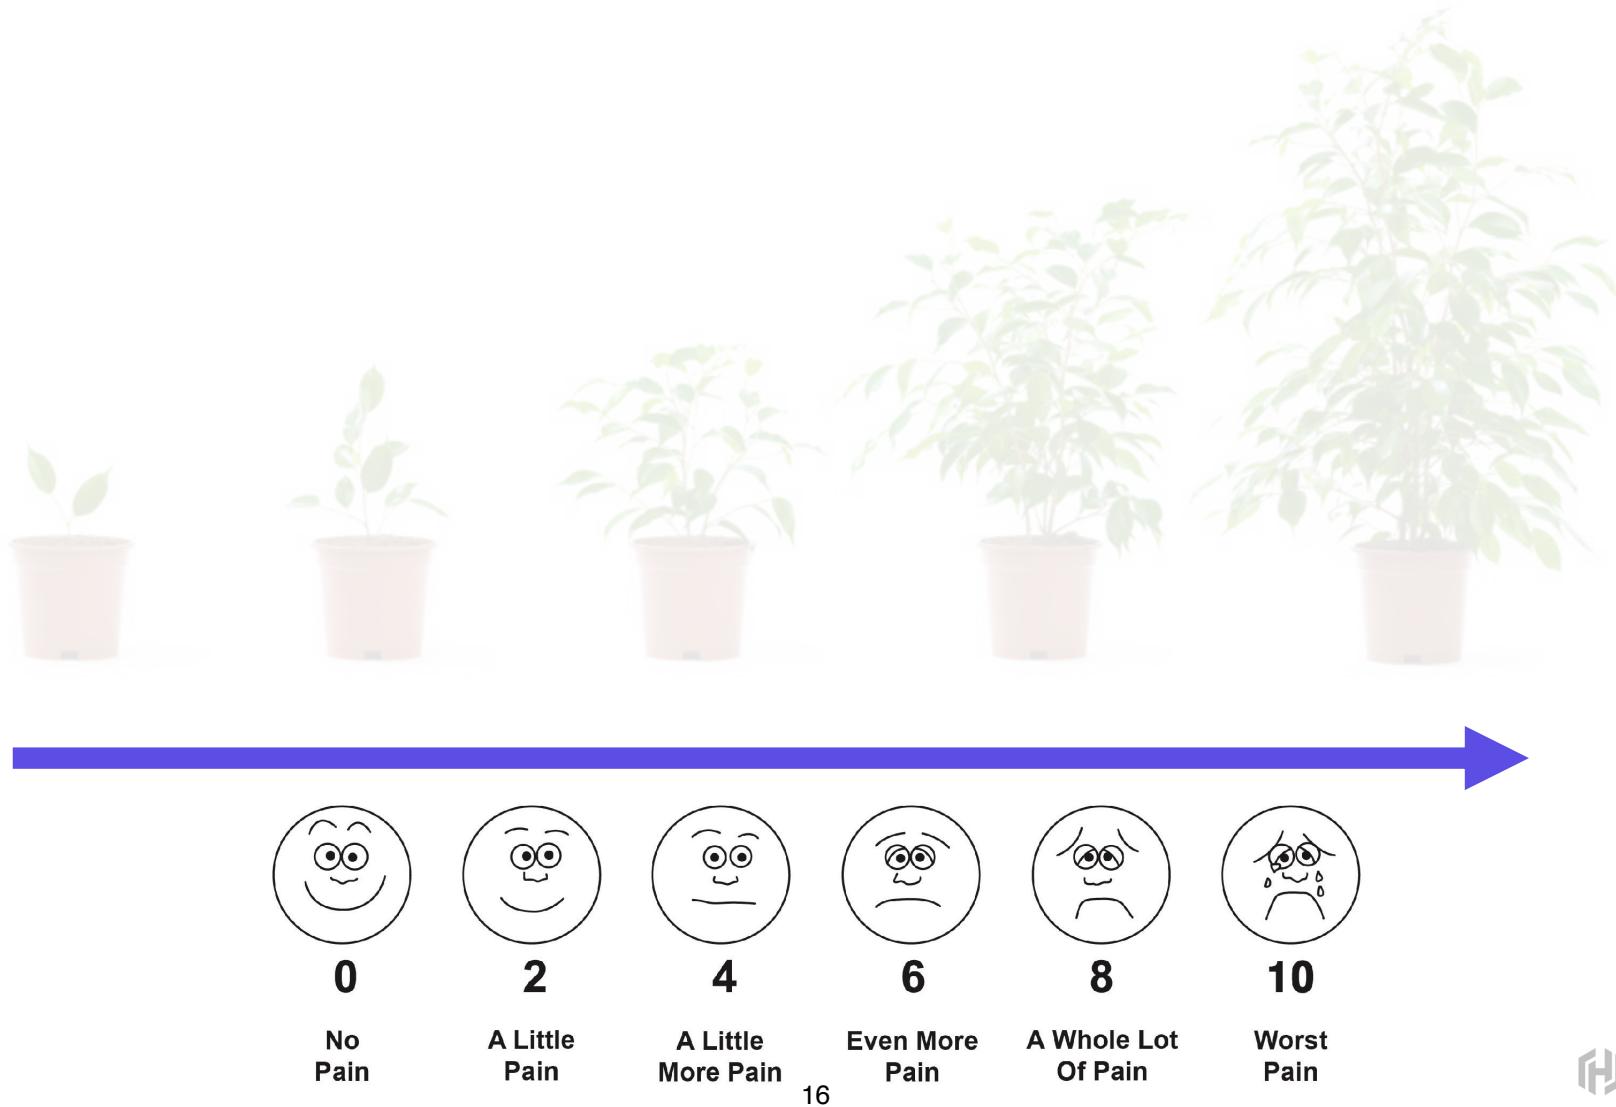
*aaS



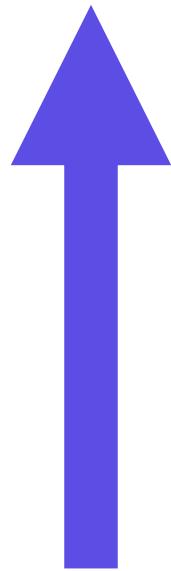
Datacenter-as-Computer



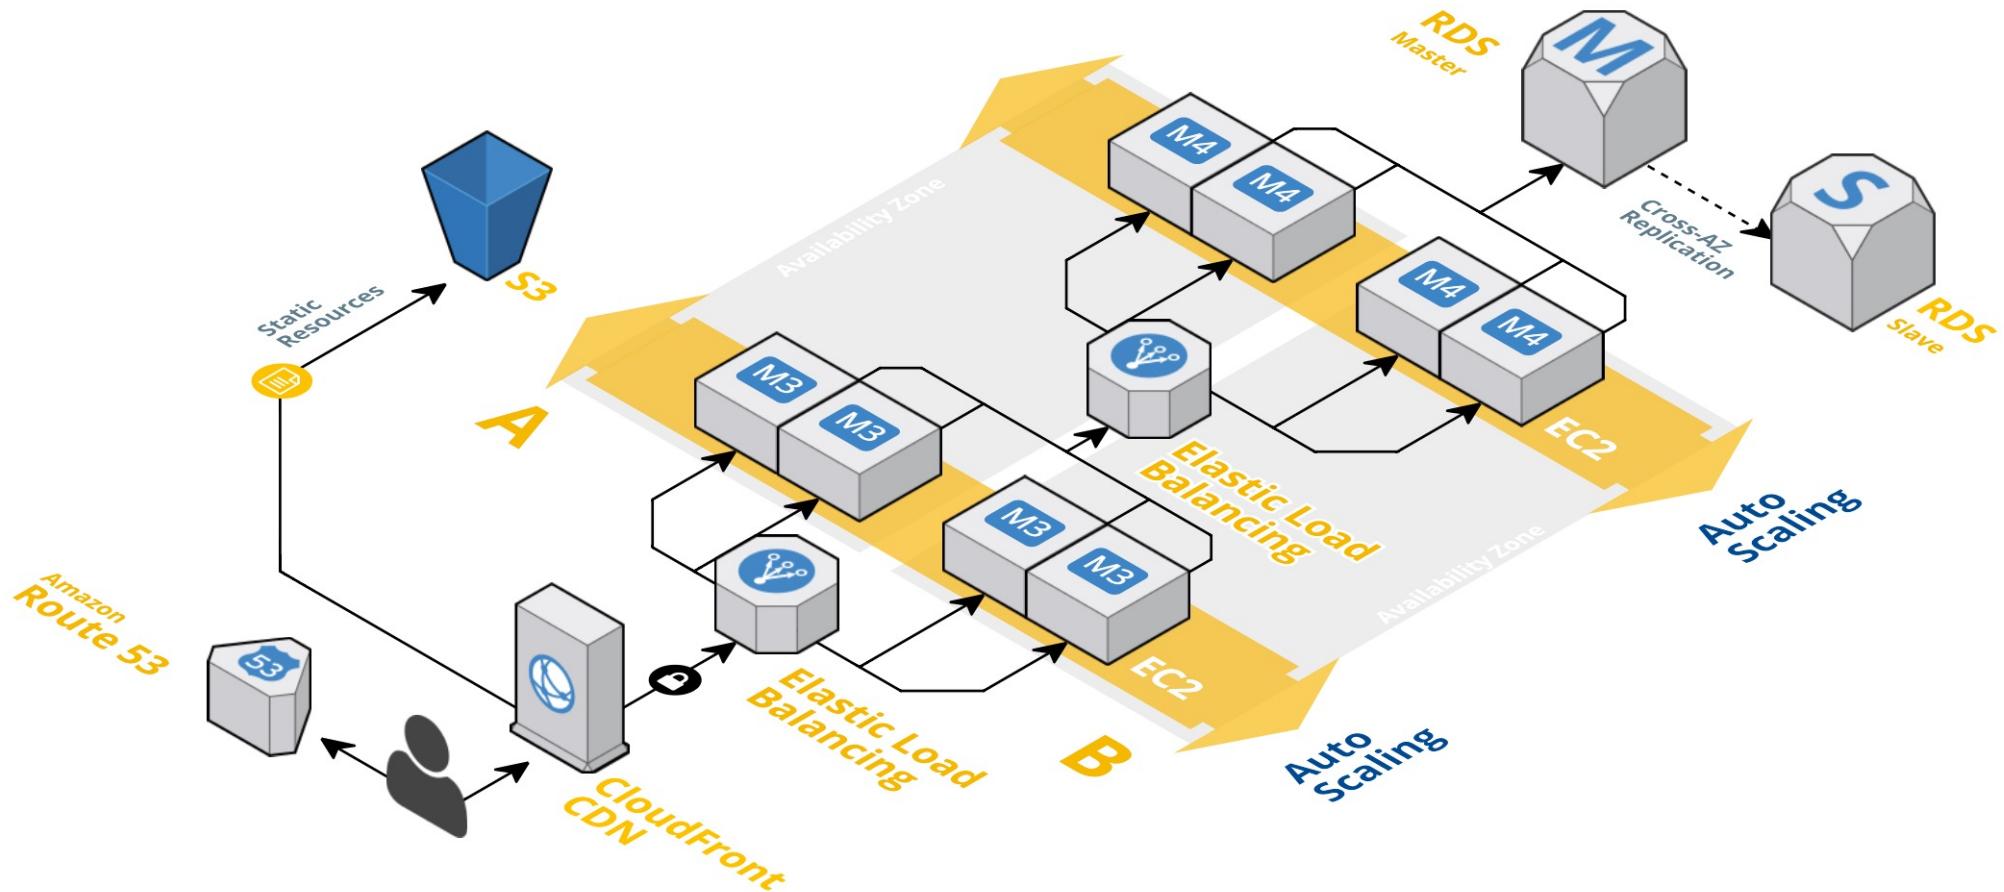
Maturity

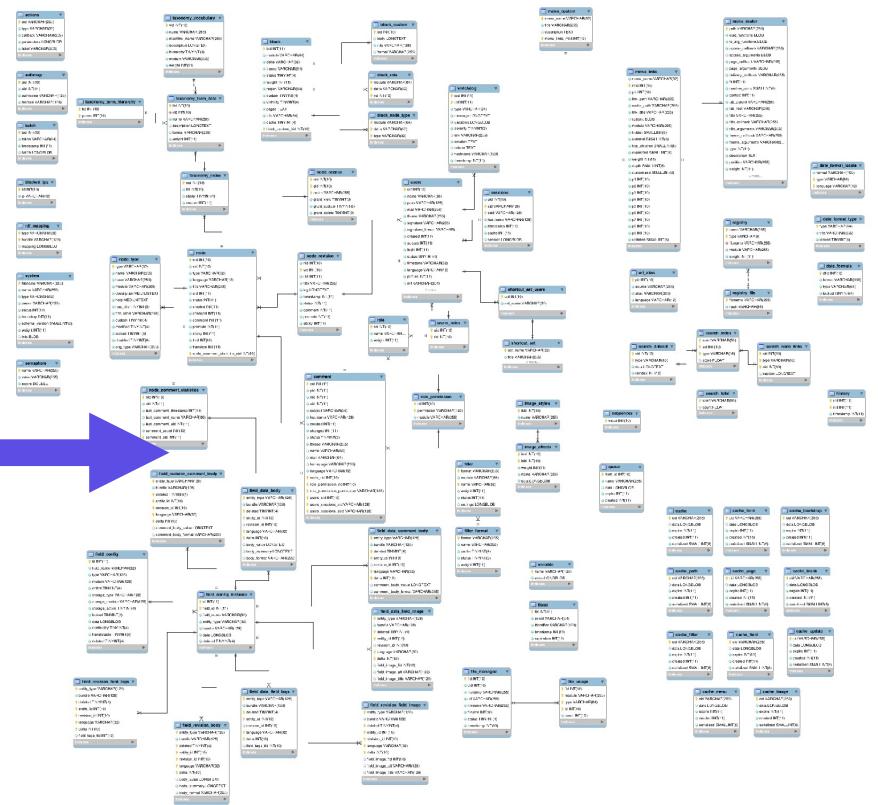
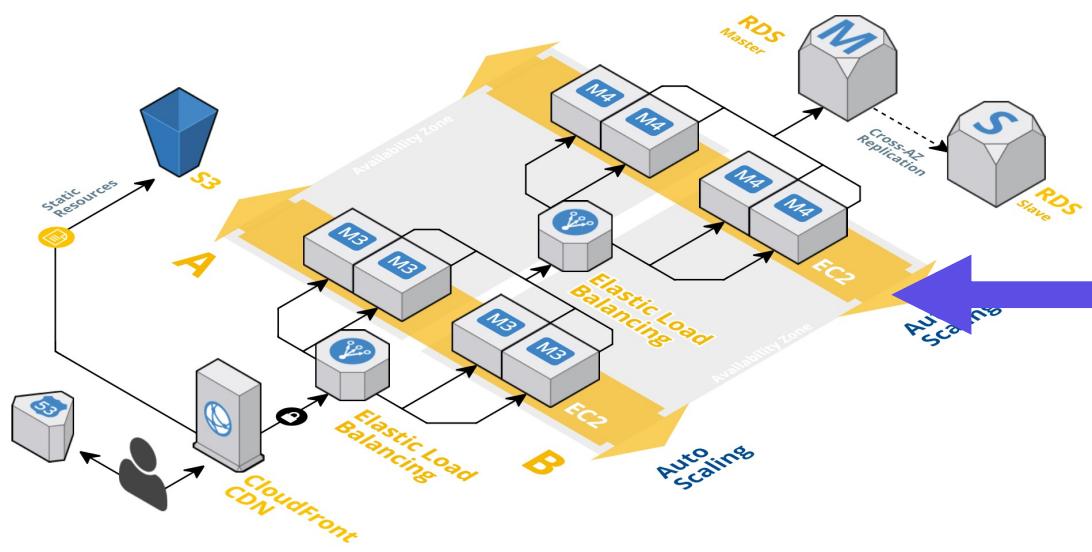


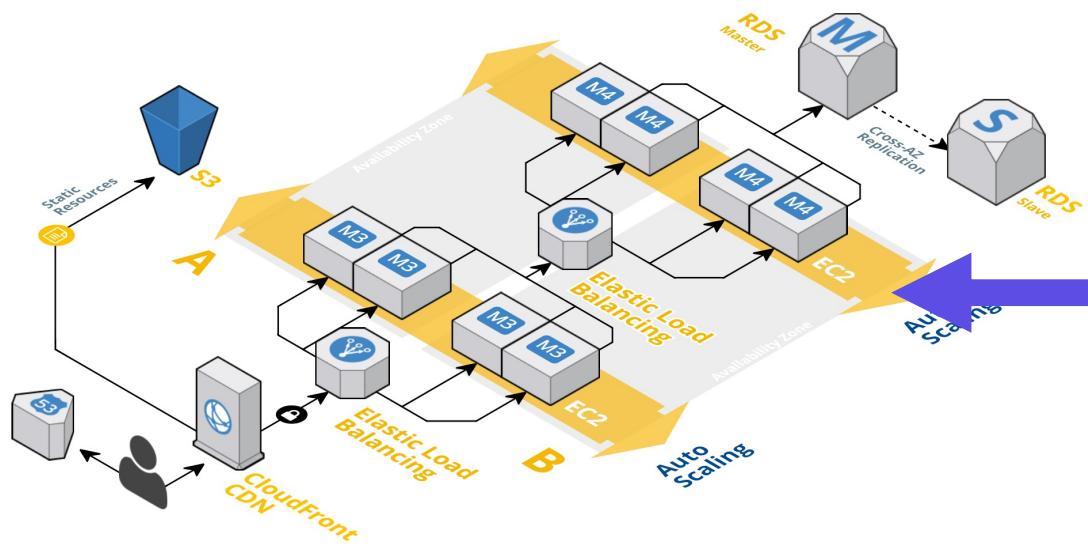
Capability



Complexity







```

475 def update
476   return update_api if api_request?
477
478   if authorized_action(@account, @current_user, :manage_account_settings)
479     respond_to do |format|
480
481       custom_help_links = params[:account].delete :custom_help_links
482       if custom_help_links
483         @account.settings[:custom_help_links] = custom_help_links.select{|k, h| h['state'] != 'deleted'}
484         hash = index_with_hash[1]
485         hash.delete('state')
486         hash.assert_valid_keys ["text", "subtext", "url", "available_to"]
487         hash
488       end
489     end
490   end
491
492   params[:account][:turnitin_host] = validated_turnitin_host(params[:account][:turnitin_host])
493   enable_user_notes = params[:account].delete :enable_user_notes
494   allow_sis_import = params[:account].delete :allow_sis_import
495   params[:account].delete :default_user_storage_quota_mb unless @account.root_account? && !@account.grants_right?(@current_user, :manage_storage_quotas)
496   [:storage_quota, :default_storage_quota, :default_user_storage_quota_mb,
497    :default_group_storage_quota, :default_group_storage_quota_mb].each { |key| params[:account].delete key }
498   end
499   if params[:account][:services]
500     params[:account][:services].slice(*Account.services_exposed_to_ui_hash(nil, @current_user, @account))
501     @account.set_service_availability(key, value == '1')
502   end
503   params[:account].delete :services
504   end
505   if @account.grants_right?(@current_user, :manage_site_settings)
506     # If the setting is present (update is called from 2 different settings forms, one for notifications)
507     # If set to default, remove the custom name so it doesn't get saved
508     params[:account][:outgoing_email_default_name] = '' if params[:account][:setting]
509   end
510
511
512   google_docs_domain = params[:account][:settings].try(:delete, :google_docs_domain)
513   if @account.feature_enabled?(:google_docs_domain_restriction) &&
514     @account.root_account? &&
515     !@account.site_admin?
516     @account.settings[:google_docs_domain] = google_docs_domain.present? ? google_docs_domain :
517   end
518
519   @account.enable_user_notes = enable_user_notes if enable_user_notes
520   @account.allow_sis_import = allow_sis_import if allow_sis_import && @account.root_account?
521   if @account.site_admin? && params[:account][:settings]
522     # these shouldn't get set for the site admin account
523     params[:account][:settings].delete(:enable_alerts)
524     params[:account][:settings].delete(:enable_eportfolios)
525   end
526   else
527     # must have :manage_site_settings to update these
528     [ :admins_can_change_passwords,
529      :admins_can_view_notifications,
530      :enable_alerts,
531      :enable_eportfolios,
532      :enable_profiles,
533      :show_scheduler,
534      :global_includes,
535      :gmail_domain
536    ].each do |key|
537      params[:account][:settings].try(:delete, key)
538

```


Terraform Concepts

Terraform's Goals

Unify the view of resources using infrastructure as code

Support the modern data center (IaaS, PaaS, SaaS)

Expose a way for individuals and teams to safely and predictably change infrastructure

Provide a workflow that is technology agnostic

Manage anything with an API

Terraform vs. Other Tools

Provides a high-level abstraction of infrastructure
(IaC)

Allows for composition and combination

Supports parallel management of resources
(graph, fast)

Separates planning from execution (dry-run)

Infrastructure as Code

Provide a codified workflow to create infrastructure

Expose a workflow for managing updates to existing infrastructure

Integrate with application code workflows (Git, SCM, Code Review)

Provide modular, sharable components for separation of concerns

Infrastructure as Code (Terraform)

Human-readable configuration (HCL) is designed for human consumption so users can quickly interpret and understand their infrastructure configuration

Terraform can also parse configuration in JSON for machine-generated configurations

Configuration format is very VCS friendly with support for multi-line lists, trailing commas, and auto-formatting



GIT DIFF EXAMPLE.TF

```
resource "azure_resource_group" "training" {
    name          = "myResourceGroup"
-   location      = "East US"
+   location      = "West US"
}
```

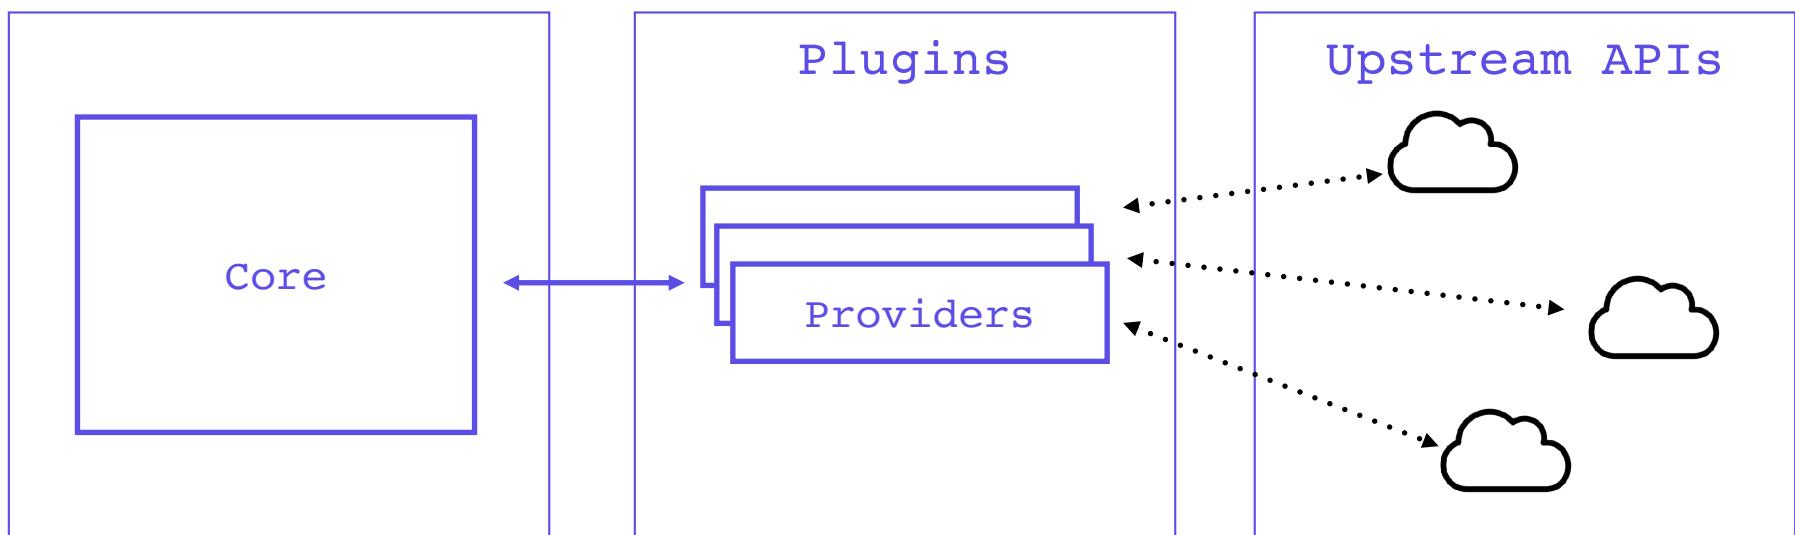
Infrastructure as Code (Terraform)

Reference values from other resources, building the implicit dependency graph.

```
resource "azurerm_network_interface" "web" {
    name          = "myWebServerNetworkInterface"
    ip_configuration= {
        private_ip_address = "10.0.2.5"
    }
}

resource "dnsimple_record" "web" {
    domain      = "hashicorp.com"
    name        = "web"
    ttl         = "3600"
    type        = "A"
    value       = azurerm_network_interface.web.private_ip_address
}
```

Terraform's Internals: Structure



Terraform's Internals: Core Concepts

Config: Target Reality

State: Current Reality

Diff: {Config – State}

Plan: Presents Diff

Apply: Resolves Diff

```
provider "azurerm" {
    features {}
}

resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}
```

Command Line Interface

All interactions with Terraform occur via the CLI

Terraform is a local tool (runs on the current machine)

The terraform ecosystem also includes providers for many cloud services, and a module repository

Hashicorp also has products to help teams manage Terraform: Terraform Cloud and Terraform Enterprise

```
> terraform help
Common commands:
apply          Builds or changes infrastructure
console          Interactive console for Terraform interpolations
destroy          Destroy Terraform-managed infrastructure
env              Workspace management
fmt              Rewrites config files to canonical format
get              Download and install modules for the configuration
graph            Create a visual graph of Terraform resources
import           Import existing infrastructure into Terraform
init           Initialize a Terraform working directory
output           Read an output from a state file
plan           Generate and show an execution plan
providers        Prints a tree of the providers used in the config
push             Upload this Terraform module to Atlas to run
refresh          Update local state file against real resources
show             Inspect Terraform state or plan
taint            Manually mark a resource for recreation
```

Terraform Init

You must run `terraform init` to initialize a new Terraform working directory, and after changing provider configuration

You **do not** need to run `terraform init` before each command

Similar to `git init`

```
> terraform init
Initializing provider plugins...
- Checking for available provider plugins on https://releases...
- Downloading plugin for provider "azurerm" (3.8)...
```

The following providers do not have any version constraints in configuration, so the latest version was installed.

To prevent automatic upgrades to new major versions that may contain breaking changes, it is recommended to add `version = "..."` constraints to the corresponding provider blocks in configuration, with the constraint strings suggested below.

```
* provider.azurerm: version = "~> 3.8"
```

```
Terraform has been successfully initialized!
```

Command: terraform plan

The plan shows you what will happen

You can save plans to guarantee what will happen

Plans show reasons for certain actions (such as re-create)

Prior to Terraform, users had to guess change ordering, parallelization, and rollout effect

Command: terraform plan

- + resource will be created
- resource will be destroyed
- ~ resource will be updated in-place
- /+ resources will be destroyed and re-created

TERMINAL

```
$ terraform plan
+ azurerm_resource_group.training
  id:                                <known after apply>
  location:                           "eastus"
  name:                               "myResourceGroup"
```

Command: `terraform apply`

Executes changes in order based on the resource graph

Parallelizes changes when possible

Handles and recovers transient errors

Runs `plan` first unless given a serialized plan file

Command: terraform apply

Updates existing resources when updates are allowed

Re-creates existing resources when updates are not allowed

```
> terraform apply
```

```
# ...
```

```
Plan: 1 to add, 0 to change, 0 to destroy.
```

```
Do you want to perform these actions?
```

```
Terraform will perform the actions described above.
```

```
Only 'yes' will be accepted to approve.
```

```
Enter a value: yes
```

```
azurerm_resource_group.training: Creating...
```

```
azurerm_resource_group.training: Still creating... [10s elapsed]
```

```
azurerm_resource_group.training: Still creating... [20s elapsed]
```

```
azurerm_resource_group.training: Still creating... [30s elapsed]
```

```
azurerm_resource_group.training: Creation complete after 31s [id=i-0b8bf14603ebdc264]
```

```
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

Command: terraform show

Displays human-friendly state

Can be used to get on-the-fly information

EXERCISE

In-Place Update

Edit the main Terraform file and add at least two tags to the Azure Resource Group.

Built-By

Environment

HINT: You may need to look at Terraform's documentation for the syntax ([terraform.io](https://www.terraform.io)).



MAIN.TF

```
provider "azurerm" {
    features {}
}

resource "azurerm_resource_group" "training" {
    name              = "myResourceGroup"
    location          = "East US"
    tags              = {
        "environment" = "production"
    }
}
```

 TERMINAL

```
$ terraform plan

~ azurerm_resource_group.training
~ tags = {
    + "environment" = "production"
}
Plan: 0 to add, 1 to change, 0 to destroy.
```

TERMINAL

```
$ terraform apply  
azurerm_resource_group.training:  
Refreshing state...
```

```
Terraform will perform the following actions:
```

```
~ azurerm_resource_group.training  
~ tags = {  
+ "environment" = "production"  
}
```

```
Apply complete! Resources: 0 added, 1 changed, 0  
destroyed.
```

TIP

Use Version Control Tools For Rollback

Terraform only knows about your configuration and the state of the infrastructure it has built.

If you need to rollback, use the capabilities of your version control software to revert to an earlier version of `main.tf`, then run `terraform apply` on it.

What You Learned

- ☑ Login to your workstation
- ☑ Create a basic Terraform configuration
- ☑ Initialize and apply the configuration (create infrastructure)
- ☑ Change and re-apply the configuration

EXERCISE

Lab 0: Instruqt Overview

Duration: 10 minutes

Goal: Familiarize yourself with your VSCode and
Instruqt Labs

EXERCISE

Lab 2 & 3: Terraform Basic Configuration and Virtual Machine

Duration: 25 minutes

Goal: Familiarize yourself with your workstation and the Terraform CLI

Terraform Configuration

Configuration Syntax

Configurations are written in HashiCorp Configuration Language (HCL)

HCL is designed to strike a balance between human-readable and machine-parsable

Terraform can also read configuration in JSON, but we recommend using HCL, except when the configuration is machine-generated

Load Order & Semantics

Configuration can be in a single file or split across multiple files.
Terraform will process all files in the current working directory
which end in `.tf` or `.tf.json`

Sub-folders are not included (non-recursive)

Files are processed in lexicographical (dictionary) order

Any files with a different extension are ignored (e.g. `.jpg` or
`.tf.old`)

Generally, the order in which things are defined doesn't matter.

Load Order & Semantics

Parsed configurations are *appended* to each other, not merged

Resources with the same name are **not** merged
(this will produce an error)

Configuration syntax is declarative, so references to other resources do not depend on the order they are defined

EXERCISE

Which Files Does Terraform Use?

Which of the following files would be parsed by Terraform when a command is issued?

```
> tree
.
├── README.md
├── main.tf
├── other.tf
├── gravatar.jpg
└── alphabet
    ├── README.md
    └── main.tf
```

```
> tree
.
├── README.md
├── main.tf
├── other.tf
├── gravatar.jpg
└── alphabet
    ├── README.md
    └── main.tf
```

Configuration Organization Patterns

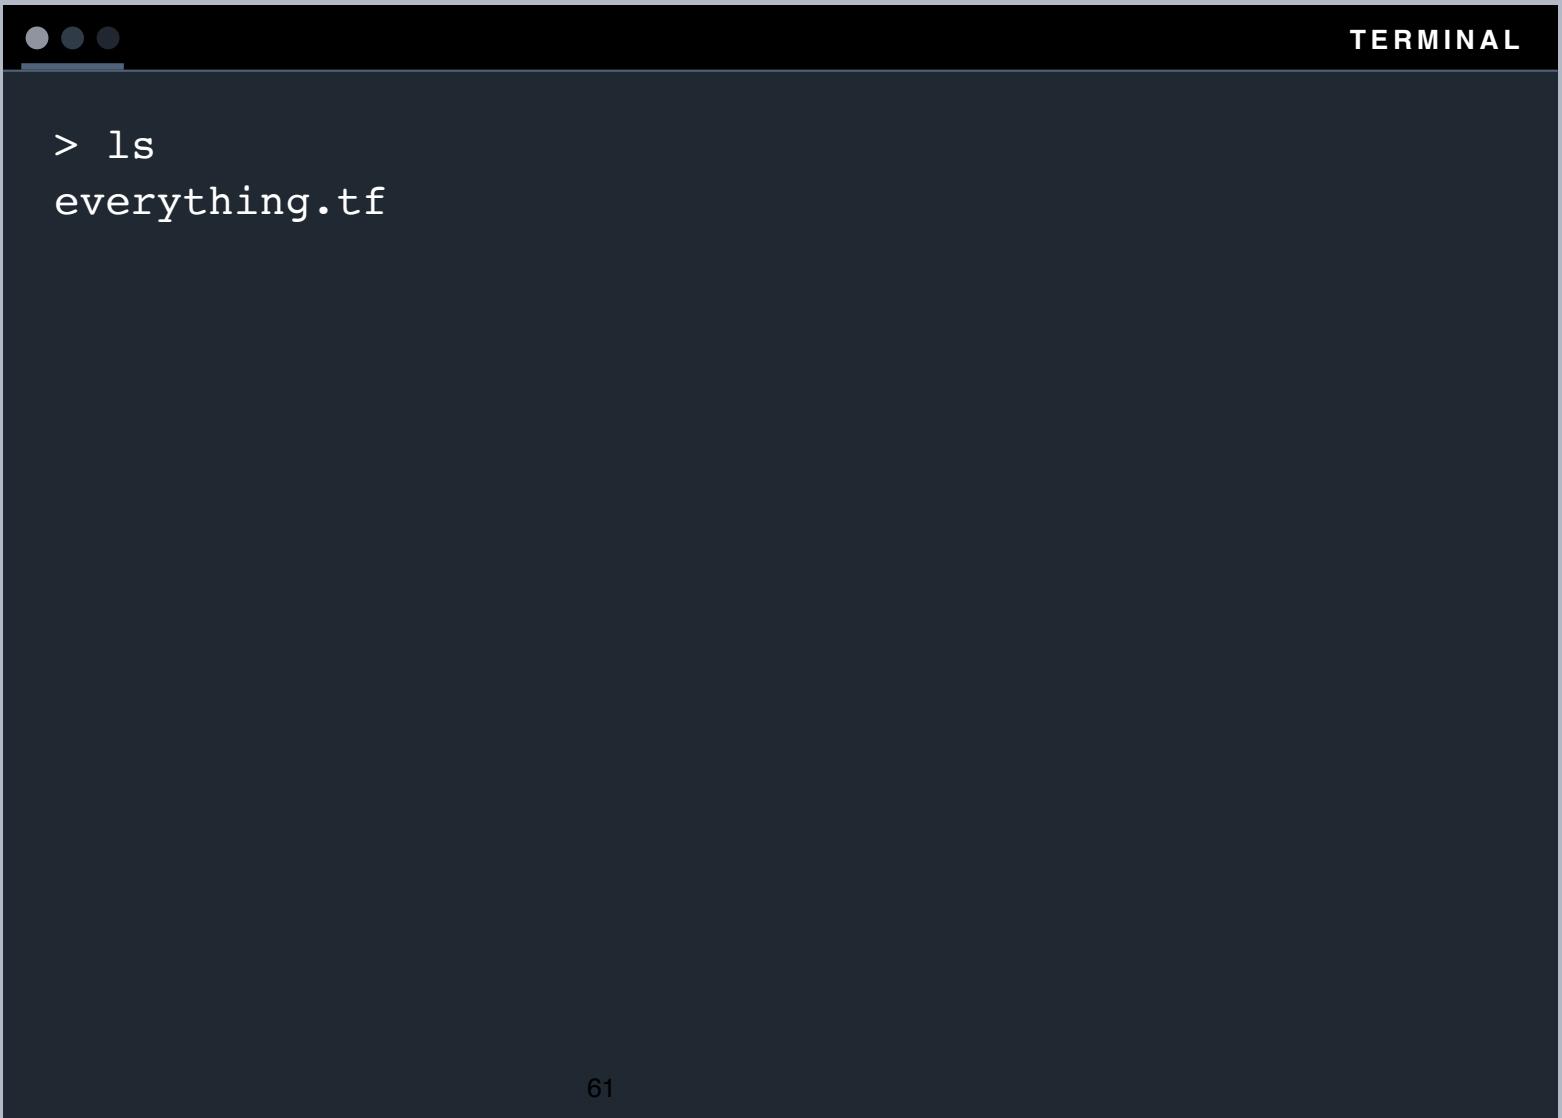
Teams choose to organize their configurations in many different patterns.

**TERMINAL**

```
> ls  
main.tf  
outputs.tf  
variables.tf
```

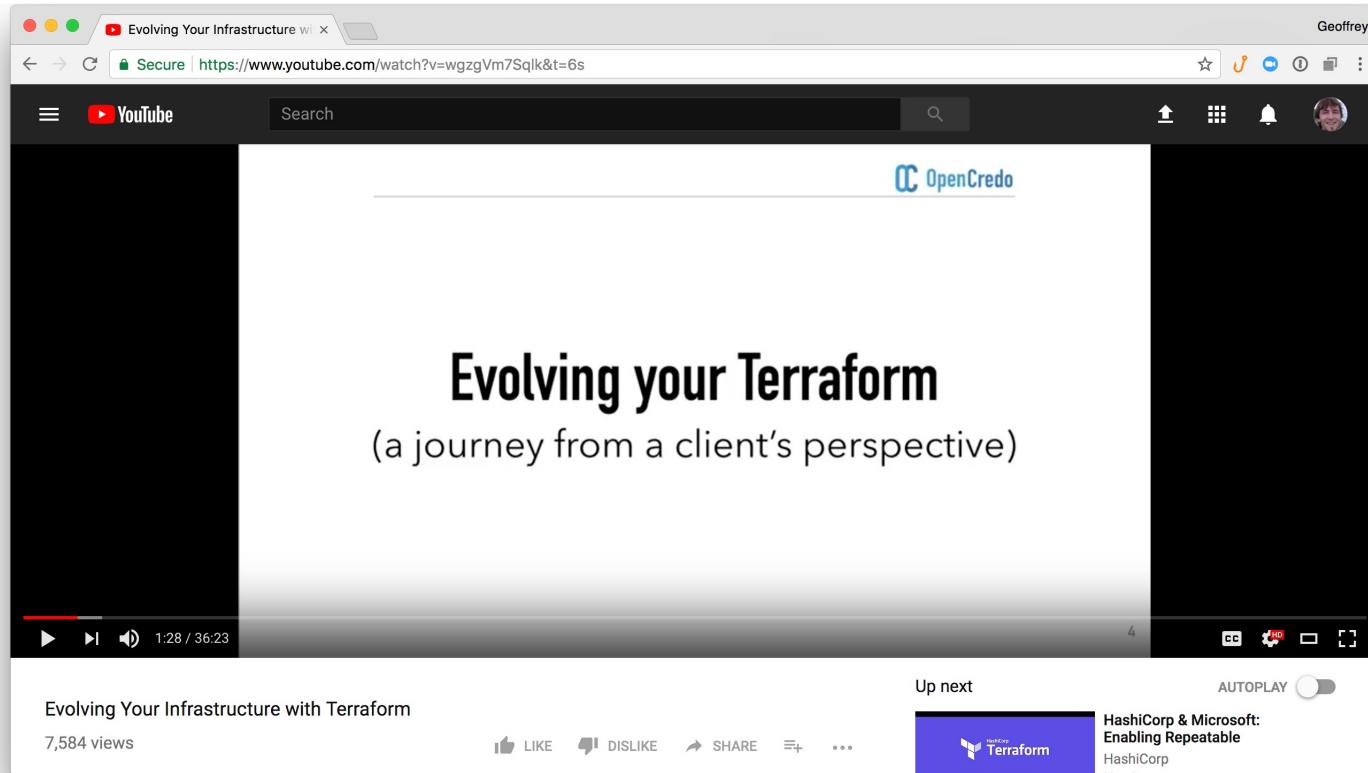
**TERMINAL**

```
> ls
instances.tf
load-balancers.tf
shared.tf
```



TERMINAL

```
> ls  
everything.tf
```



```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}
```

Top Level Keywords

provider

variable

resource

output

module

data

terraform

locals

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}
```

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}
```

MAIN.TF

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = production
    }
}

resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = production
    }
}
```

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = production
    }
}

resource "azurerm_virtual_machine" "web-2" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = production
    }
}
```

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = production
    }
}

resource "digitalocean_droplet" "web" {
    image    = "ubuntu-18-04-x64"
    name     = "web"
    size     = "s-1vcpu-1gb"
    region   = "nyc2"
}
```



MAIN.TF

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size                = "Standard_F2"
    tags                   = {
        "environment" = "production"
    }
}
```

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}
```

```
resource "azurerm_network_interface" "web" {
    name          = "myWebServerNetworkInterface"
    ip_configuration = {
        private_ip_address = "10.0.2.5"
    }
}

resource "dnsimple_record" "web" {
    domain      = "hashicorp.com"
    name        = "web"
    ttl         = "3600"
    type        = "A"
    value       = azurerm_network_interface.web.private_ip_address
}
```

```
# This is a comment

resource "azurerm_network_interface" "web" {
    name          = "myWebServerNetworkInterface"
    ip_configuration = {
        private_ip_address = "10.0.2.5"
    }
}

/* This is a multiline
comment */

resource "dnsimple_record" "web" {
    # ...
}
```

```
variable "thing" {
    description = <<EOF
This is a
multiline
string
EOF
}
```

Syntax Highlighting

Plugins for Terraform/HCL exist for most major editors, but ruby tends to work best if one does not exist.

Outputs

What You'll Learn

- 🕒 Print dynamic data with an output
- 🕒 Query specific values from the output

Outputs

Outputs define useful values that will be highlighted to the user when Terraform applies: IP addresses, usernames, generated keys

Outputs can be queried using the `terraform output` command

Outputs are a way to easily extract and query information from all of Terraform's collected data

```
resource "azurerm_resource_group" "training" {
# ...
}

output "location" {
    value = azurerm_resource_group.training.location
}

# Prior to Terraform 0.12 syntax:

output "location" {
    value = "${azurerm_resource_group.training.location}"
}
```

```
resource "azurerm_resource_group" "training" {
# ...
}

output "location" {
    value = azurerm_resource_group.training.location
}

output "name" {
    value = azurerm_resource_group.training.name
}
```

```
> terraform refresh
azurerm_resource_group.training: Refreshing state... [id=i-
0b8bf14603ebdc264]
```

Outputs:

```
location = "eastus"
name      = "myResourceGroup"
```

 TERMINAL

```
> terraform output
location = "eastus"
name      = "myResourceGroup"
```

TERMINAL

```
> terraform output location  
"eastus"
```

What You Learned

- 🕒 Print dynamic data with an output
- 🕒 Query specific values from the output

EXERCISE

Lab 4: Outputs

Duration: 10 minutes

Goal: Configure specific outputs and learn to query your configuration for those.

Console

What You'll Learn

- 🕒 Interact with live data in the console
- 🕒 Use the console from the command line

Command: `terraform console`

The `terraform console` command creates an interactive console for interacting with Terraform

This is useful for experimenting with interpolations and interacting with Terraform state

 TERMINAL

```
> terraform console
> azurerm_resource_group.training.location
  eastus
> azurerm_resource_group.training.name
  myResourceGroup
> azurerm_resource_group.training.tags
{
  "environment" = "production"
}
> <Ctrl-C>
```

 TERMINAL

```
> terraform console  
> exit
```

 TERMINAL

```
> echo "azurerm_resource_group.training.location" |  
terraform console  
  
"eastus"
```

Functions

Terraform has built-in functions to achieve common data manipulation tasks

Functions categories include strings, collections, encoding, and encryption.

Functions can be tested using the `terraform console` command



MAIN.TF

```
resource "azurerm_virtual_machine" "web" {
    name                  = "my_server"
    location              = "East US"
    resource_group_name   = "my_rg"
    network_interface_ids = ["azure_network_interface.main.id"]
    vm_size               = "Standard_F2"
    tags                  = {
        "environment" = "production"
    }
}

output "server_name" {
    value = lower(azurerm_virtual_machine.web.name)
}
```

EXERCISE

Lab 5: Console

Duration: 10 minutes

Goal: The Read-Evaluate-Print-Loop is useful for understanding interpolation and finding specific information.

Variables

What You'll Learn

- 🕒 Use variables to pass values to your configuration
- 🕒 Refactor existing configuration to use variables
- 🕒 Keep sensitive data out of source code control
- 🕒 Pass variables to Terraform in several ways

Variables

Work a lot like variables from programming languages

Allow you to remove hard coded values and pass parameters to your configuration

Can help make configuration easier to understand and re-use

Must always have a value. Variables are never optional, but they can have a default value

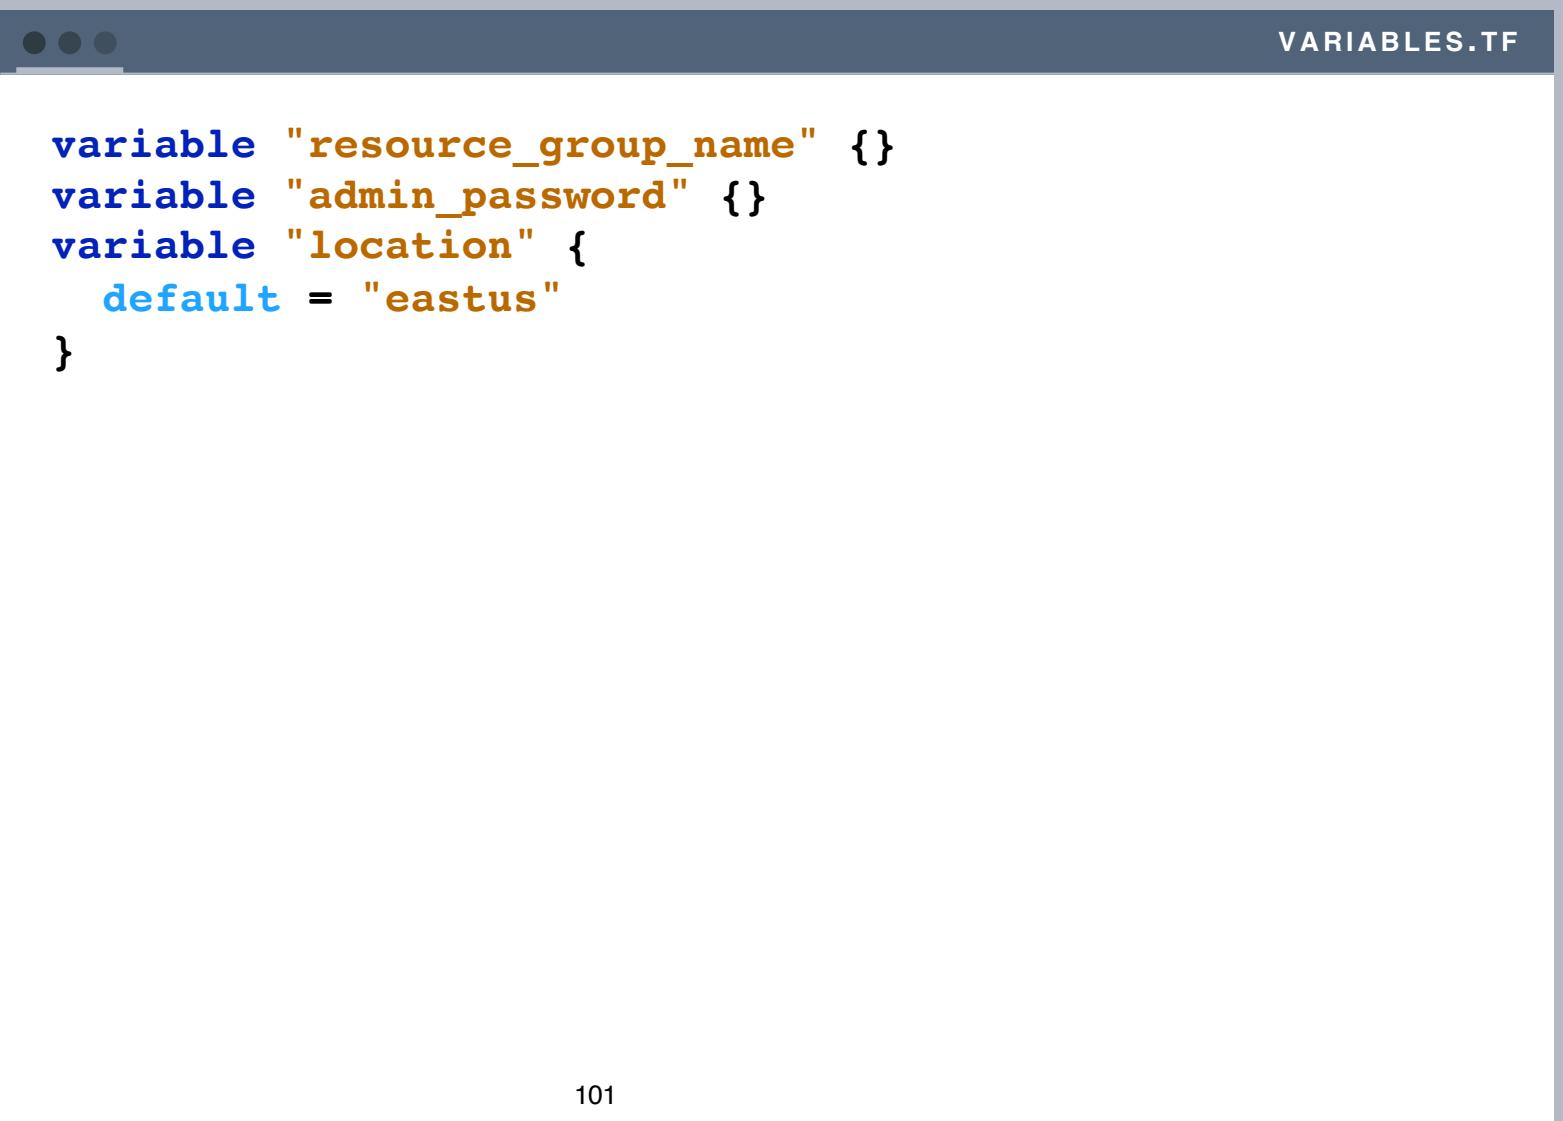
Variable Types

The most common variable types are `string`, `number`, `list`, and `map`

Other support types include `bool` (true/false), `set`, `object`, and `tuple`

If omitted, the type is inferred from the default value

If neither type nor default is provided, the type is assumed to be `string`



VARIABLES.TF

```
variable "resource_group_name" {}
variable "admin_password" {}
variable "location" {
    default = "eastus"
}
```

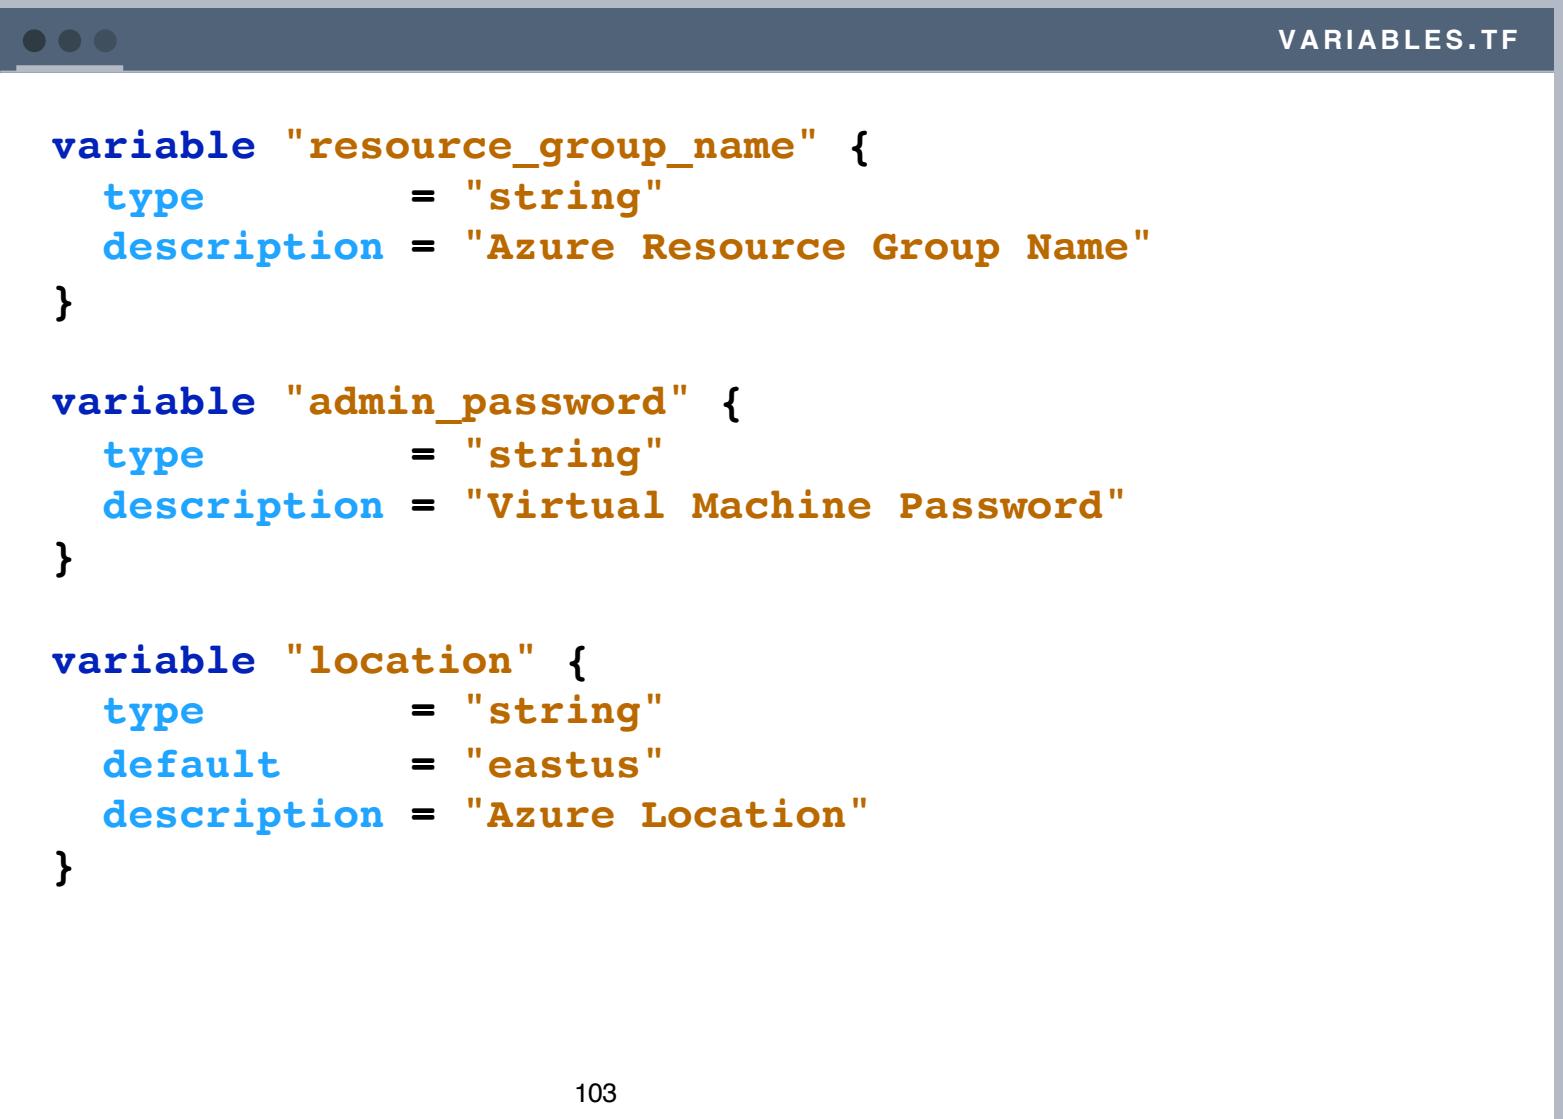


VARIABLES.TF

```
variable "resource_group_name" {
    type = "string"
}

variable "admin_password" {
    type = "string"
}

variable "location" {
    type     = "string"
    default = "eastus"
}
```



VARIABLES.TF

```
variable "resource_group_name" {
    type      = "string"
    description = "Azure Resource Group Name"
}

variable "admin_password" {
    type      = "string"
    description = "Virtual Machine Password"
}

variable "location" {
    type      = "string"
    default   = "eastus"
    description = "Azure Location"
}
```

Using Variables

Variables can be referred to by using the syntax
var.<name>

Variables' values can also be interpolated into a string using the syntax “ ***\${var.<name>}*** ”

```
resource "azurerm_resource_group" "training" {
    name      = var.resource_group_name
    location  = var.location
    tags      = {
        "environment" = "production"
    }
}
```

● ● ● TERMINAL

```
> terraform plan  
var.resource_group_name  
Enter a value:
```

```
<Ctrl + C>
```

● ● ● TERMINAL

```
> export TF_VAR_resource_group_name=<rg_name>  
  
> terraform plan  
var.admin_password  
Enter a value:  
  
<Ctrl + C>
```

A screenshot of a code editor window titled "TERRAFORM.TFVARS". The window contains a block of Terraform configuration code. The code defines variables for a resource group, admin password, location, environment tag, prefix, computer name, and admin username.

TERRAFORM.TFVARS

```
resource_group_name  = "MyResourceGroupName"
admin_password       = "Password1234!"
location             = "East US"
# EnvironmentTag      = "staging"
# prefix               = "appA"
# computer_name        = "myserver"
# admin_username        = "testadmin"
```

 TERMINAL

```
> terraform plan -var admin_password=<secret_key>  
azurerm_virtual_machine.web: Refreshing state... (ID: i-  
02f5717f1a84502ed)
```

```
# ...
```

```
No changes. Infrastructure is up-to-date.
```

```
# ...
```

TERMINAL

```
> terraform plan -var-file secrets.tfvars
```

Specifying Variable Values

1. Default values
2. Environment variables
3. Saved in variable definition files (`terraform.tfvars`)
4. Using the `-var` or `-var-file` command line option
5. Configured in a Terraform Enterprise or Terraform Cloud workspace
6. Entered via CLI prompts



VARIABLES.TF

```
variable "num_webservers" {
    type    = "number"
    default = 5
}

variable "regions" {
    type    = "list"
    default = ["eastus", "westus"]
}

variable "image_ids" {
    type    = "map"
    default = {
        us-east-1 = "image-1234"
        us-west-2 = "image-4567"
    }
}
```

WARNING

Variables & Security

Never put secret values, like passwords or access tokens, in `.tf` files or other files that are checked into source control.

Locals

Locals are values defined inside a Terraform configuration.

Locals can construct values from variables, resources, data source, and other locals.

Locals are scoped to their module.



VARIABLES.TF

```
locals {
    naming_prefix = "ant"
    default_name  = "${local.naming_prefix}-${var.company}"
    number_list   = [5,7,9]
}
```

What You Learned

- ☑ Use variables to store sensitive values
- ☑ Refactor existing configuration to use variables
- ☑ Keep sensitive data out of source code control systems
- ☑ Populate variables to Terraform in several ways

EXERCISE

Lab 6: Variables

Duration: 15 minutes

Goal: Remove all hardcoded values in your current configuration and use variables instead.

Auto-Formatting/File Layout

Auto-Formatting

Terraform has built-in support for auto-formatting configuration files to match the HCL specification.

There is no need to manually align things.

This helps keep configurations VCS-friendly and reduces bike-shed arguments over formatting styles.

Plugins exist for most major editors.

Command: `terraform fmt`

The `terraform fmt` command formats configuration files.

It does not prompt for confirmation or

It also reports on syntax errors, if they exist.

EXERCISE

Format Your Code

Run the `terraform fmt` command with no arguments. By default it will choose all Terraform configurations in the current working directory non-recursively.

Open the `main.tf` file and see that it has been formatted.

Command: terraform fmt

Please feel free to use the `terraform fmt` command through this training and at your leisure, but we will not explicitly discuss it anymore.

Command: `terraform validate`

The `terraform validate` checks the syntax and logic of your configuration

Checks both core terraform and provider resources

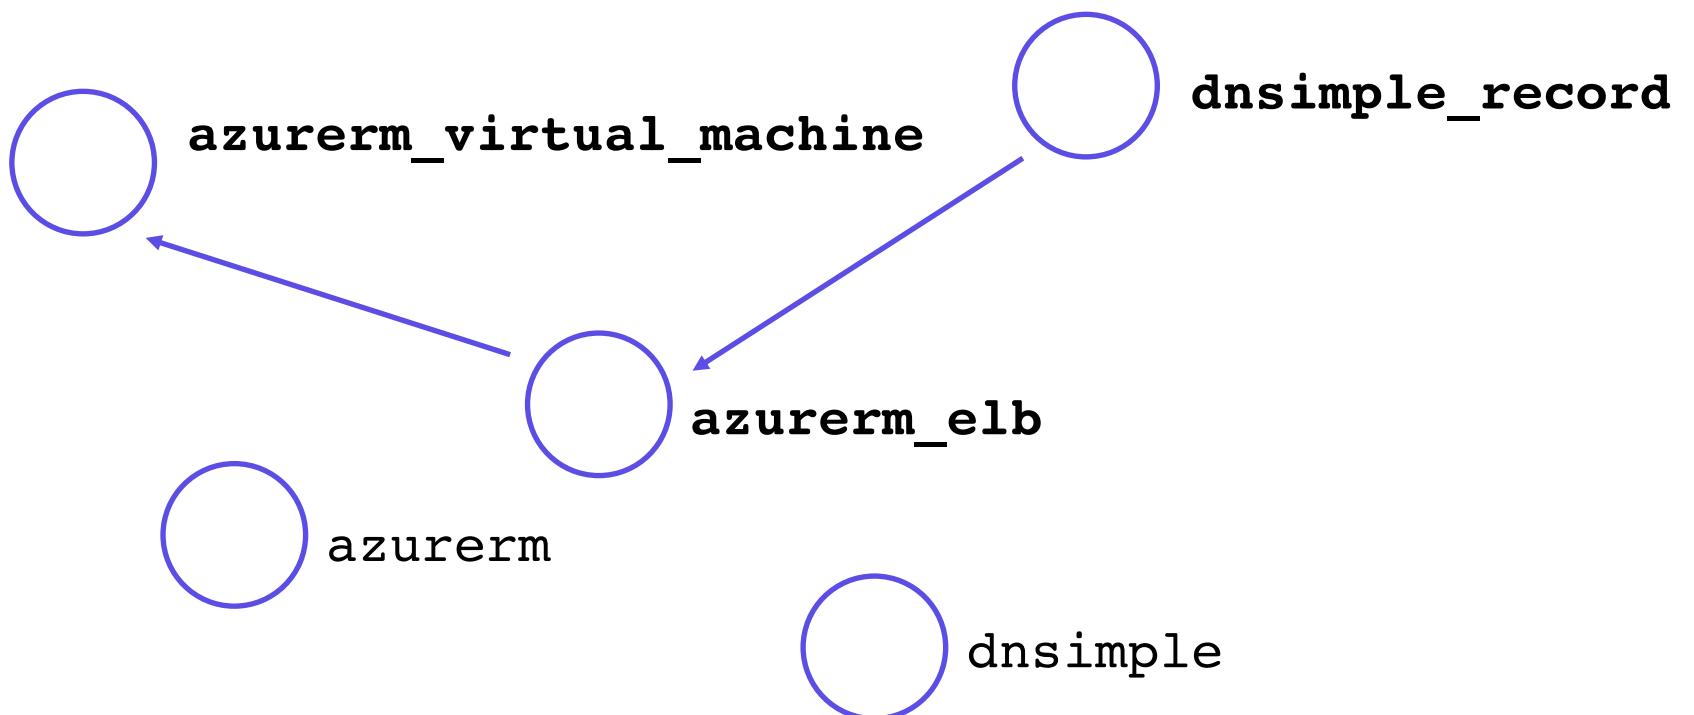
Requires initialization before running

Graph Theory

What You'll Learn

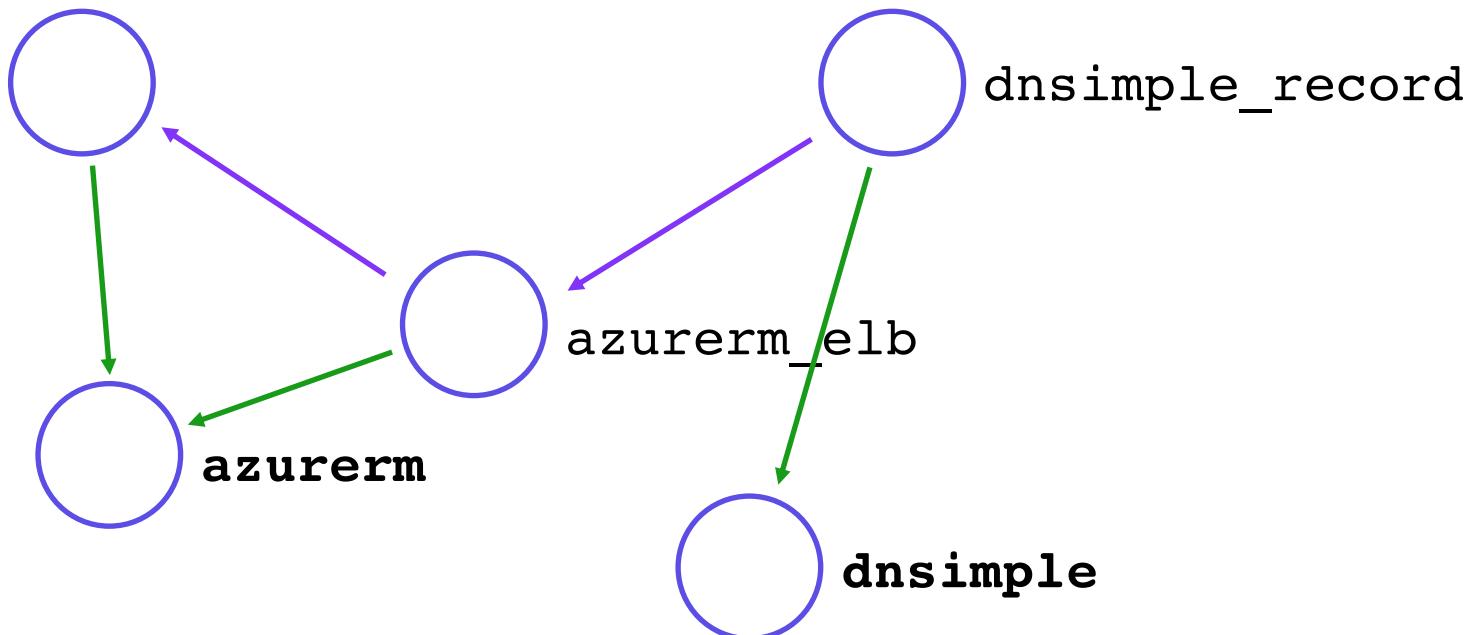
- ☑ Understand how graph theory makes Terraform more effective
- ☑ Generate a graph of your infrastructure

Building the Graph



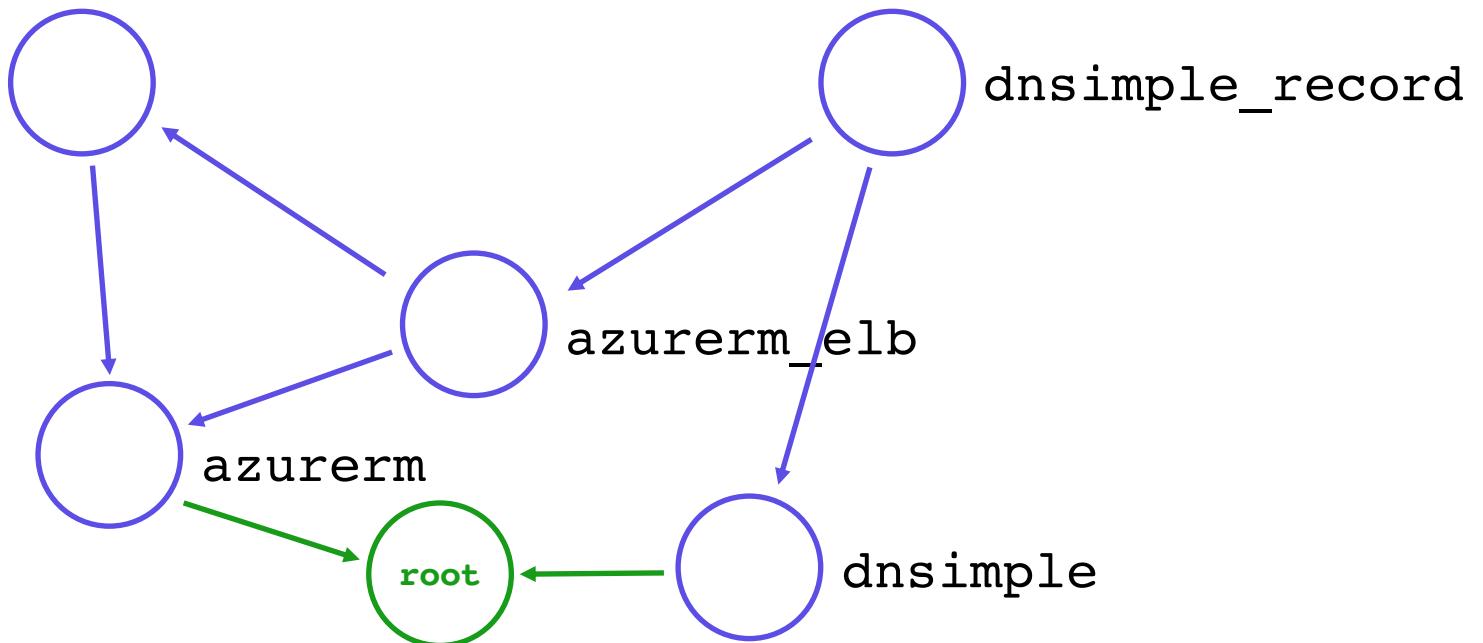
Building the Graph

`azurerm_virtual_machine`



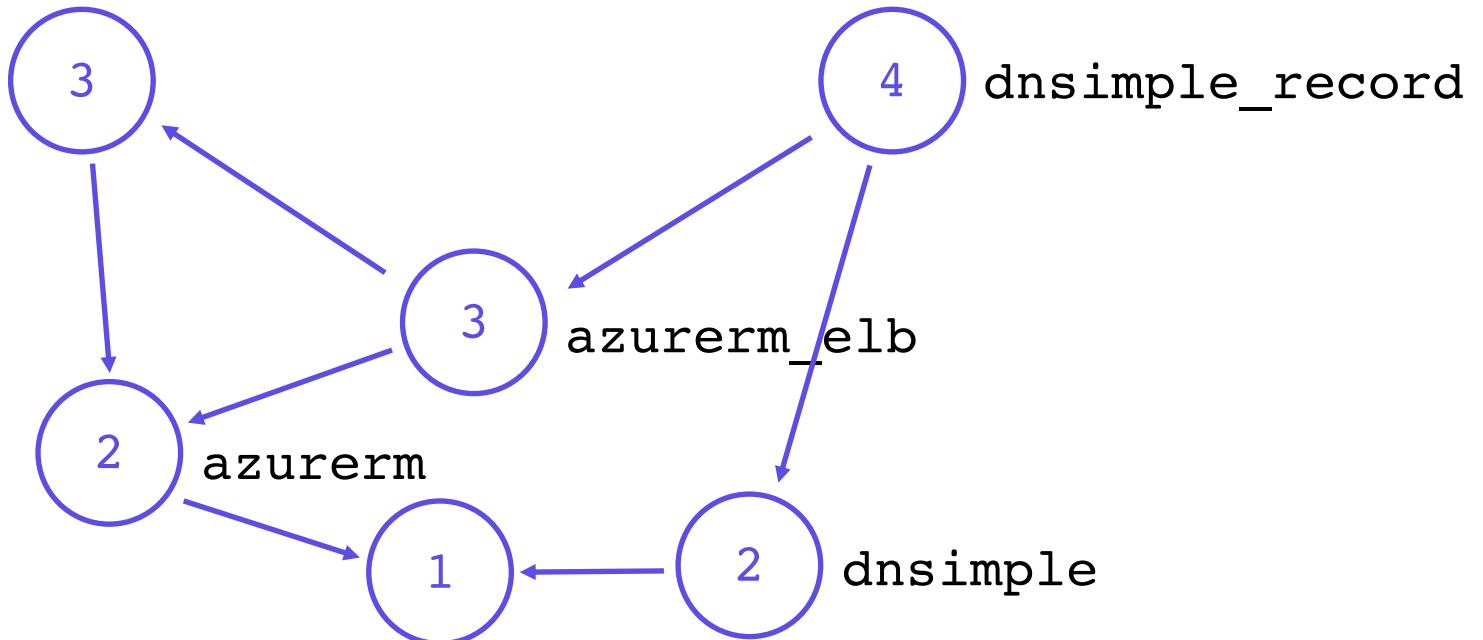
Building the Graph

azurerm_virtual_machine



Walking the Graph

azurerm_virtual_machine



Resource Graph

The resource graph is an internal representation of all resources and their dependencies.

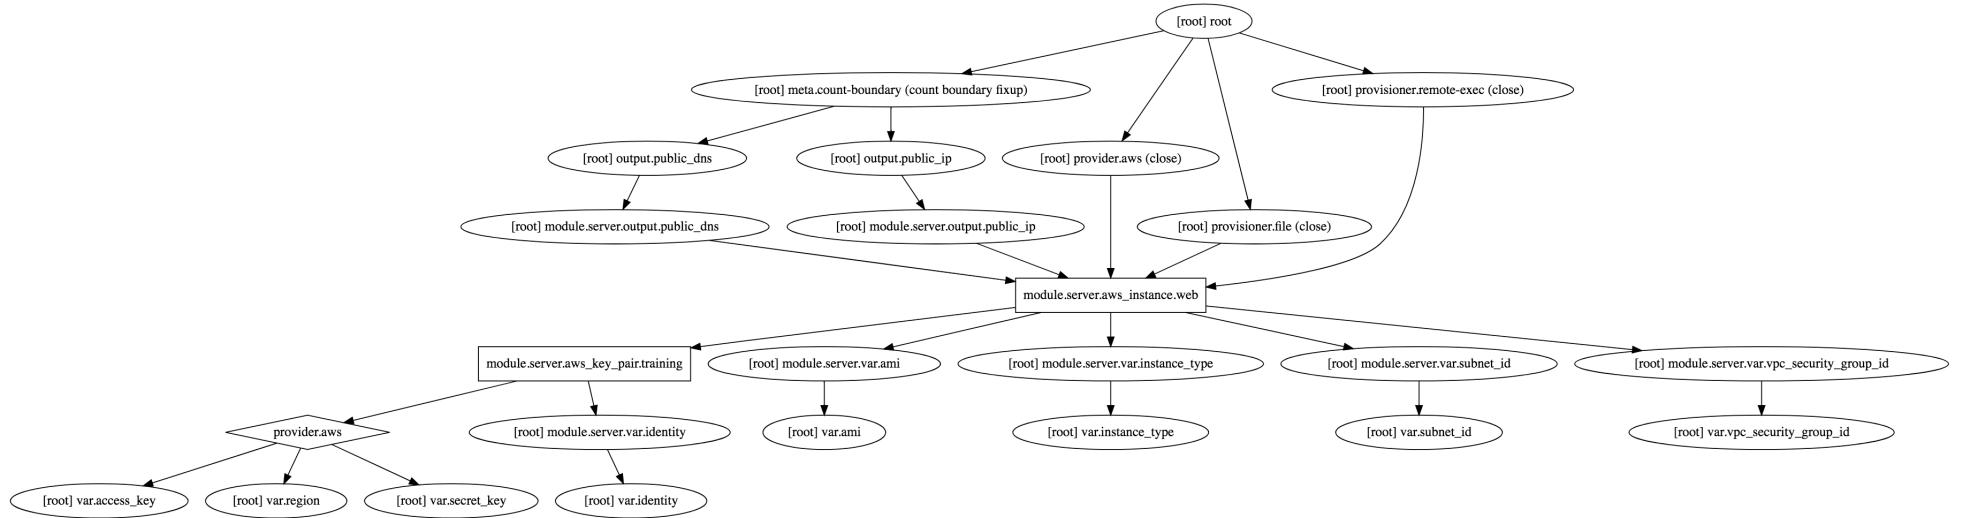
A human-readable graph can be generated using the `terraform graph` command.

Can optionally draw cycles (advanced).

 TERMINAL

```
> terraform graph
digraph {
    compound = "true"
    newrank = "true"
    subgraph "root" {
        "[root] azurerm_virtual_machine.web" [label =
"azurerm_virtual_machine.web", shape = "box"]
        "[root] provider.azurerm" [label = "provider.azurerm", shape = "diamond"]
        "[root] azurerm_virtual_machine.web" -> "[root] provider.azurerm"
        "[root] meta.count-boundary (count boundary fixup)" -> "[root] ..."
        "[root] provider.azurerm (close)" -> "[root] azurerm_virtual_machine.web"
        "[root] root" -> "[root] meta.count-boundary (count boundary fixup)"
        "[root] root" -> "[root] provider.azurerm (close)"
    }
}
```

terraform graph



Terraform Graph

Useful for visualizing infrastructure and dependencies

Builds upon existing visualization technologies and open formats such as DOT

Can optionally draw cycles (resources that depend on each other in a circle)

EXERCISE

Lab 7: Graphs

Duration: 10 minutes

Goal: Visualize the dependencies in a Terraform configuration with the graph command.

Meta Arguments

What You'll Learn

- 🕒 Understand meta arguments
- 🕒 Use count to create multiple instances
- 🕒 Manage your own metadata

Meta-Arguments

Meta-arguments allow for higher-level control flow and lifecycle management in Terraform

These don't map directly to cloud resources or APIs, but help you control Terraform's actions

Meta-Arguments

Count

The count argument allows for N number of identical resources to be created. This removes the need for iteration with "for" or "while" loops in many cases

Meta-Arguments

For-Each

The `for_each` argument allows for N number of resources to be created based on a map or set. Similar to `count`, the `for_each` argument gives greater flexibility over each created resource.

Meta-Arguments

Depends On

The `depends_on` argument allows for declaration of explicit dependencies. This is useful where interpolation is not required, but explicit ordering is desired

Meta-Arguments

Provider

The `provider` argument allows for a resource to explicitly use a provider by name or alias. This is most useful when there are multiple providers in a single configuration, such as `aws.west` and `aws.east`

Meta-Arguments

Lifecycle

The `lifecycle` argument allow explicit configuration of resource lifecycle such as preventing destruction or ignoring property changes. This is an advanced option and is not recommended for users who are getting started with Terraform



SERVER/SERVER.TF

```
# ...

resource "azurerm_virtual_machine" "web" {
    count          = 2

    name           = "myServer-${count.index}"
    location       = "eastus"

# ...
}
```

 SERVER/SERVER.TF

```
# ...

output "public_ip" {
    value = azurerm_network_interface.web.*.public_ip
}

output "public_dns" {
    value = azurerm_network_interface.web.*.public_dns
}
```

EXERCISE

Lab 8: Meta Arguments

Duration: 10 minutes

Goal: Modify the number instances created with count.

```
> terraform apply
# module.server.azurerm_virtual_machine.web[1] will be
  created
+ resource "azurerm_virtual_machine" "web" {
    + name:                  "myServer-1"
    + location:              "useast"
# ...

Plan: 1 to add, 0 to change, 0 to destroy.

# ...

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```



SERVER/SERVER.TF

```
variable "num_webservers" {
    count    = number
    default  = 2
}

# ...

resource "azurerm_virtual_machine" "web" {
    count = var.num_webservers

    name          = "myServer-${count.index}"
    location      = "eastus"

    tags {
        # ...
        "Name" = "web ${count.index+1}/${var.num_webservers}"
    }
}
```

Provisioners

What You'll Learn

- ⌚ How to create a connection to our server
- ⌚ Specify the key for connecting to the server
- ⌚ Provision the server with files and commands

WARNING

The Server We Created is Useless

We created a server without any running code.

No useful services are running on it.

In this section, we'll see how to provision or instance.

```
resource "azurerm_virtual_machine" "web" {
    # . . .

    connection = {
        type      = "ssh"
        user      = "ubuntu"
        host      = self.public_ip
        private_key = module.keypair.private_key_pem
    }
}
```



SERVER/MAIN.TF

```
resource "azurerm_virtual_machine" "web" {
    # . . .

    provisioner "file" {
        source      = "assets"
        destination = "/tmp/"
    }

}
```

<https://www.terraform.io/docs/provisioners/file.html>



SERVER/MAIN.TF

```
resource "azure_virtual_machine" "web" {
    # . . .

    provisioner "remote-exec" {
        inline = [
            "sudo sh /tmp/assets/setup-web.sh"
        ]
    }
}
```

<https://www.terraform.io/docs/provisioners/remote-exec.html>

TIP

Force Provisioning With the `-replace` argument

If we run `apply` after adding provisioners, it won't change the server.

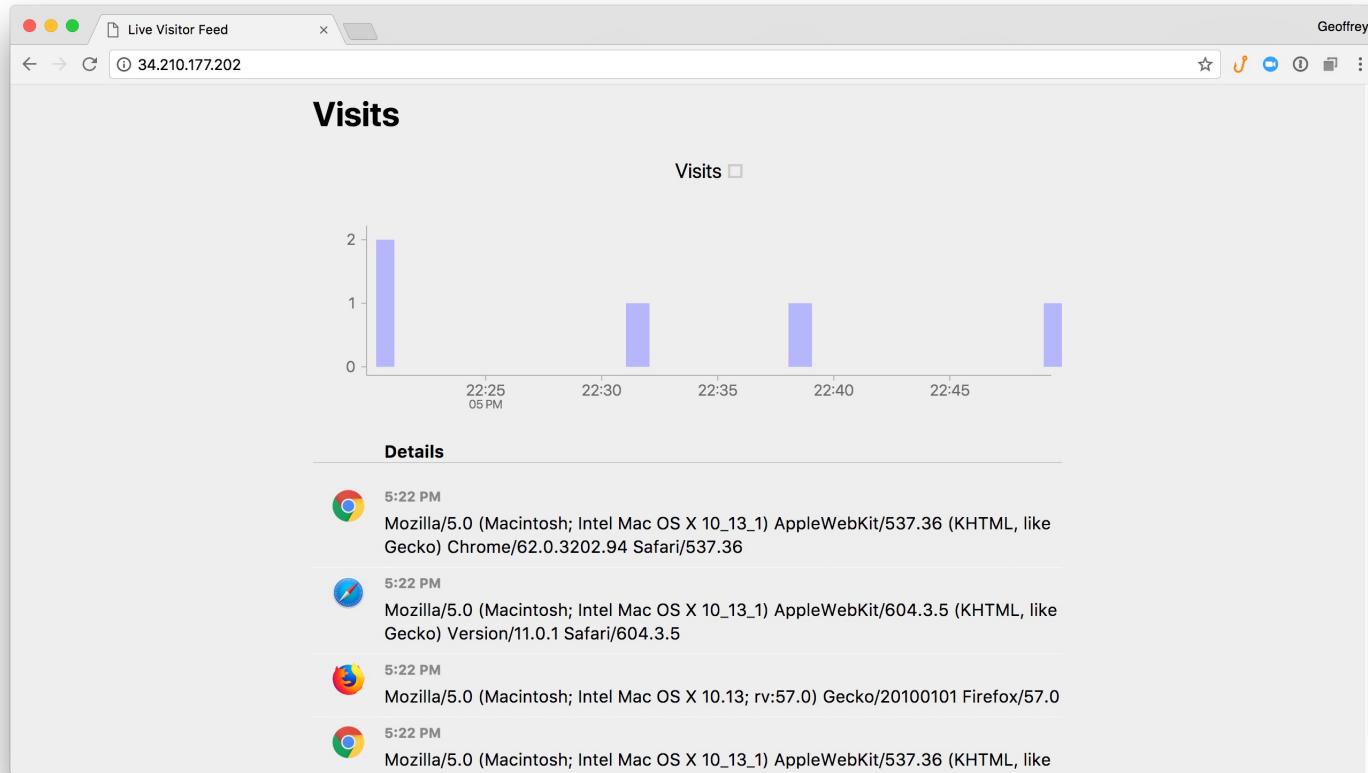
Force it to be destroyed and recreated with the `-replace` argument:

```
terraform apply -replace azurerm_virtual_machine.web
```

If you're unsure of the correct ID for the resource you want to replace, you can use `terraform show`:

```
> terraform show
azurerm_virtual_machine.web:
  id = "/subscriptions/9c3c5800-dc5b-4723-8b57-
205b827d6a06/resourceGroups/myResourceGroup/providers/Microsoft.C
ompute/virtualMachines/web_vm"
  name = myWebServer
  . . .
```

```
> terraform apply -replace azurerm_virtual_machine.web  
# ...  
Plan: 1 to add, 0 to change, 1 to destroy.  
Do you want to perform these actions?  
Terraform will perform the actions described above.  
Only 'yes' will be accepted to approve.  
  
Enter a value: yes  
  
# ...  
  
web: Provisioning with 'remote-exec'...  
web (remote-exec): Connecting to remote host via SSH...  
web (remote-exec): Host: 35.162.171.124  
web (remote-exec): User: ubuntu  
web (remote-exec): Password: false  
web (remote-exec): Private key: true  
web (remote-exec): SSH Agent: false  
web (remote-exec): Connected!
```



Provisioning Best Practices

Good

Use `remote-exec` provisioner on a base AMI to run a few commands on instance creation.

Better

Automate the creation of infrastructure and initialization of instances with `user_data` or AWS/Azure cloud-init.

Best

Build AMIs with Packer so that minimal configuration is needed.

What You Learned

- 🕒 How to create a connection to our server
- 🕒 Specify the key for connecting to the server
- 🕒 Provision the server with files and commands

EXERCISE

Lab 9: Provisioners

Duration: 10 minutes

Goal: Provisioners allow us to create instances that are ready to use at runtime. Use your SSH key from the last lab to connect to and install our web app on a new instance.

Modules

What You'll Learn

- 🕒 The Terraform Public Module Registry
- 🕒 Use a module from `main.tf`
- 🕒 Understand how modules are stored on disk

Modules

Portable Terraform configurations (packages)

Allow separation of concerns and responsibilities
among teams

Parallels: Chef Cookbook, Puppet Module, Ruby
gem

Modules

Modules are just Terraform configurations inside a folder - there's nothing special about them.

Complex configurations, team projects, and multi-repository codebases will benefit from modules.
Get into the habit of using them wherever it makes sense.

 TERMINAL

```
> tree my-module
my-module
└── main.tf
```

```
variable "location" {}
variable "name" {}

resource "azurerm_network_interface" "web" {
    name          = "myWebServerNetworkInterface"
    ip_configuration = {
        private_ip_address = "10.0.2.5"
    }
}

output "ip_address" {
    value =
    azurerm_network_interface.web.ipconfiguration.private_ip_address
}
```

```
module "my-module" {
  # Source can be any URL or file path
  source = "../../my-module"

}
```

```
module "my-module" {
    # Source can be any URL or file path
    source  = "../../my-module"
    location = "eastus"
    name = "myServer"
}
```

```
module "my-module" {
    # Source can be any URL or file path
    source  = "../../my-module"
    location = "eastus"
    name      = "myServer"
}

output "ip-address" {
    value      = module.my-module.ip_address
}
```

Modules

Variables defined in a module become arguments to module block

Outputs defined in module become attributes from module definition

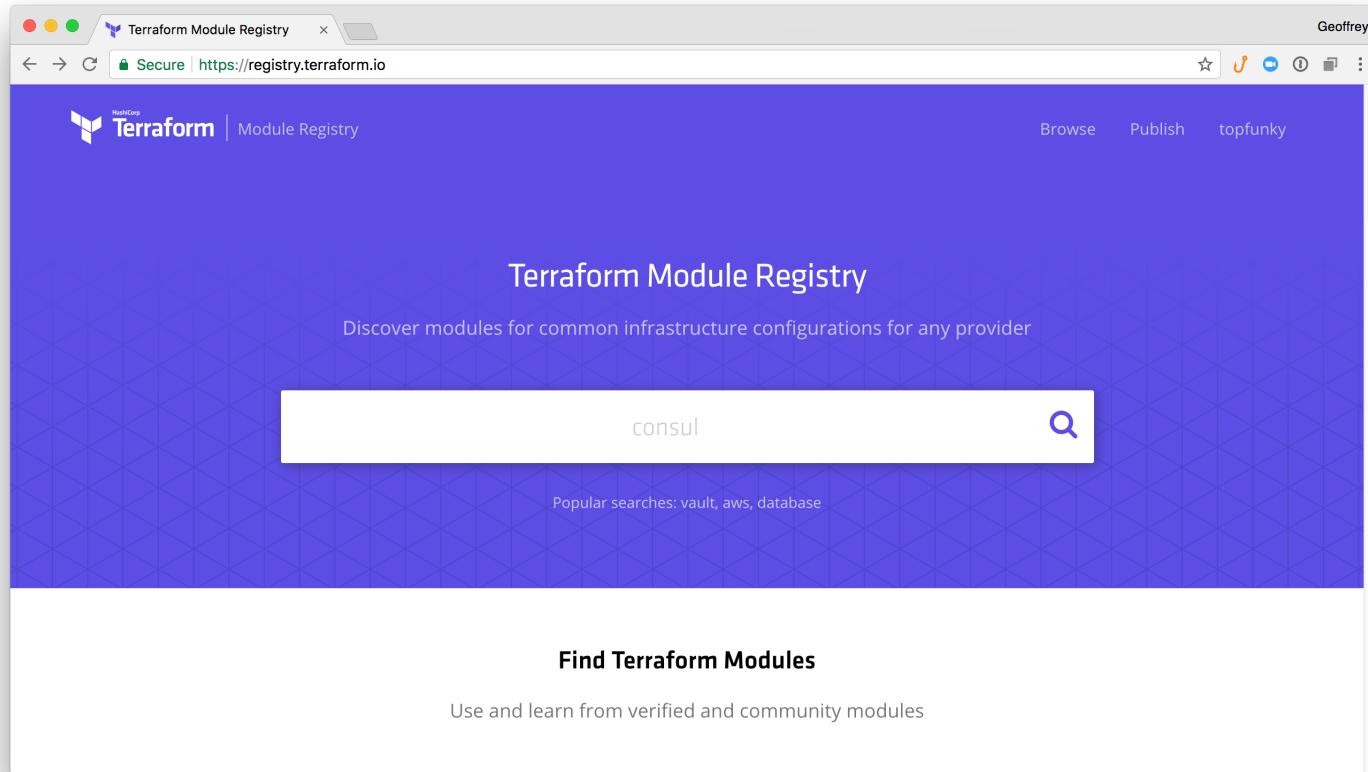
Configuration inside a module is a "black box" from the outside

Individual resources are not accessible outside the module

```
module "my-module" {
  # Source can be any URL or file path
  source = "../../my-module"

  argument_1 = "value"
  argument_2 = "value"
}

output "example" {
  value = module.my-module.azure_network.web.public_ip
}
```



Terraform Registry

Search Providers and Modules

Browse ▾ Publish ▾ Sign-in

network 

AZURERM

Terraform Azure RM Module for Network

Published May 31, 2021 by Azure

Module managed by [yupwei68](#)

Source Code: github.com/Azure/terraform-azurerm-network ([report an issue](#))

Version 3.5.0 (latest) ▾

Module Downloads All versions ▾

Downloads this week	2,642
Downloads this month	12,292
Downloads this year	44,005
Downloads over all time	177,783

Provision Instructions

Copy and paste into your Terraform configuration, insert the variables, and run `terraform init`:

```
module "network" {  
    source  = "Azure/network/azurerm"  
    version = "3.5.0"  
    # insert the 1 required variable  
}
```



173

```
module "network" {
    source          = "Azure/network/azurerm"
    version         = "3.5.0"
    resource_group_name = azurerm_resource_group.name
}
```

```
> terraform plan
Failed to load root config module: Error loading modules:
module example: not found, may need to be downloaded
using 'terraform get'
```

```
> terraform plan
```

```
Error: Could not satisfy plugin requirements
```

```
# ...
```

```
Error: provider.tls: no suitable version installed  
version requirements: "(any version)"  
versions installed: none
```

```
Error: provider.local: no suitable version installed  
version requirements: "(any version)"  
versions installed: none
```

```
> terraform init
Initializing modules...

Initializing the backend...

Initializing provider plugins...
- Checking for available provider plugins...
- Downloading plugin for provider "null" (terraform-providers/null) 2.1.2...
- Downloading plugin for provider "tls" (terraform-providers/tls) 2.0.1...
- Downloading plugin for provider "local" (terraform-providers/local) 1.3.0...

# ...
```

TIP

Module References Are Stored in `.terraform`

A hidden directory is built at the root of your project

You can use the `tree` command or `ls -l` to see the contents

Note that local modules are symlinked. This explains why changes to module code are available immediately

TERMINAL

```
> tree .terraform

.terraform/
└── modules
    └── 56e6098163e5b2675d14437a2a723a54
        └── terraform-azurerm-network-37d94b3
            ├── main.tf
            ├── outputs.tf
            ├── README.md
            └── variables.tf
    └── modules.json
└── plugins
    └── linux_amd64
        ├── lock.json
        ├── terraform-provider-azurerm_v1.23.0_x4
        ├── terraform-provider-local_v1.1.0_x4
        ├── terraform-provider-null_v1.0.0_x4
        └── terraform-provider-tls_v1.1.0_x4

4 directories, 10 files
```

 TERMINAL

```
> terraform apply
azurerm_resource_group.training: Refreshing state... [id=i-
03724b60e0ff853bc]

# ...

# module.network.azurerm_network_interface.training will be created
+ resource "azurerm_resource_group" "training" {
    + id          = (known after apply)
    + resource_group_name = "myResourceGroup"
}

# ...

Plan: 5 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes
```

What You Learned

- ☑ Import a module from the Terraform Public Module Registry
- ☑ Use a module from `main.tf`
- ☑ Understand how modules are stored on disk

EXERCISE

Lab 11: Modules

Duration: 30 minutes

Goal: Refactor your existing code into a separate module and connect a module from the Public Registry.

Destroy

What You'll Learn

- Destroy infrastructure that Terraform has created

Command: `terraform destroy`

Destroys running infrastructure

Does not touch infrastructure not managed by Terraform

```
> terraform destroy
azurerm_resource_group.training: Refreshing state...
[ id=i-0ec4ffe55abd300a5 ]
```

An execution plan has been generated and is shown below. Resource actions are indicated with the following symbols:

- destroy

Terraform will perform the following actions:

```
# azurerm_resource_group.training will be destroyed
- resource "azurerm_resource_group" "training" {
# ...
```

EXERCISE

Lab 10: Destroy

Duration: 5 minutes

Goal: Use the destroy command to destroy your infrastructure

Course Agenda

Terraform Fundamentals

Students will walk away with a solid understanding of HashiCorp Terraform.

Course Agenda:

- . Day 1 Recap
- . Data Sources
- . Terraform State
- . Examining Terraform State
- . Import
- . Terraform Cloud Basics

Course Agenda

Terraform Fundamentals

Students will walk away with a solid understanding of HashiCorp Terraform.

Course Agenda:

- . Lifecycles
- . Template Files and Rendering
- . Debug
- . Certification Prep

Azure Authentication

Authentication to Azure

Azure CLI

`az login`

- Uses your own credentials
- Great for local development
- Also how the Azure Cloud Shell works



Authentication to Azure

Service Principal

- Preferred for shared environments
- Works well in automation
- Also how Terraform Enterprise works



Authentication to Azure

Service Principal

```
provider "azurerm" {  
  
    subscription_id = "SUBSCRIPTION-ID"  
    client_id      = "CLIENT-ID"  
    client_secret  = "CLIENT-SECRET"  
    tenant_id      = "TENANT-ID"  
}
```



Authentication to Azure

Service Principal

```
export ARM_TENANT_ID=
export ARM_SUBSCRIPTION_ID=
export ARM_CLIENT_ID=
export ARM_CLIENT_SECRET=
```



Authentication to Azure

Managed Service Identity (MSI)

<no configuration needed>



Azure Location

```
resource "azurerm_resource_group" "main" {  
    location = "eastus"  
}
```

```
resource "azurerm_resource_group" "main" {  
    location = "East US"  
}
```



Data Sources

Data Sources

Data sources let you query information defined outside of Terraform.

Similar to a resource, data sources are part of a provider.

Data sources are read-only and refreshed during a plan or apply.

```
data "azurerm_subnet" "web_subnet" {
    name          = "web"
    virtual_network_name = "data-source-network"
    resource_group_name  = "data-source-network"
}

# ...

resource "azurerm_network_interface" "web" {
    name          = "web-nic"
    location      = "eastus"
    resource_group_name = data.azurerm_subnet.web_subnet.resource_group_name

    ip_configuration {
        name          = "internal"
        subnet_id     = data.azurerm_subnet.web_subnet
        .id
        private_ip_address_allocation = "Dynamic"
    }
}
```

EXERCISE

Lab 2: Data Sources

Duration: 10 minutes

Goal: Use a data source for the subnet ID of a network interface

State

What You'll Learn

- 🕒 Understand how state is managed
- 🕒 Understand where state can be stored
- 🕒 Understand Terraform Cloud and Terraform Enterprise

State

Terraform stores the state of your managed infrastructure from the last time Terraform was run.

Terraform uses this state to create plans and make changes to your infrastructure.

It is critical that this state is maintained appropriately so future runs operate as expected.

```
> head terraform.tfstate
{
  "version": 4,
  "terraform_version": "1.0.7",
  "serial": 3,
  "lineage": "067aac2b-fc72-0bad-92a4-bd48547e28e5",
  "outputs": {},
  "resources": [
    {
      "mode": "managed",
      "type": "azurerm_resource_group",
```

TERMINAL

```
> terraform state list
```

```
azurerm_resource_group.data_source
azurerm_virtual_network.data_source
```

State Resolutions

Configuration	State	Reality	Operation
azurerm_virtual_machine.web			
azurerm_virtual_machine.web	azurerm_virtual_machine.web		
azurerm_virtual_machine.web	azurerm_virtual_machine.web	azurerm_virtual_machine.web	
	azurerm_virtual_machine.web	azurerm_virtual_machine.web	
		azurerm_virtual_machine.web	
azurerm_virtual_machine.web		azurerm_virtual_machine.web	
	azurerm_virtual_machine.web		

State Resolutions

Configuration	State	Reality	Operation
azurerm_virtual_machine.web			create
azurerm_virtual_machine.web	azurerm_virtual_machine.web		create
azurerm_virtual_machine.web	azurerm_virtual_machine.web	azurerm_virtual_machine.web	noop
	azurerm_virtual_machine.web	azurerm_virtual_machine.web	delete
		azurerm_virtual_machine.web	noop
azurerm_virtual_machine.web		azurerm_virtual_machine.web	re-create*
	azurerm_virtual_machine.web		update state

State Locking

If supported, the state backend will "lock" to prevent concurrent modifications which could cause corruption.

Not all backends support locking - Terraform's documentation identifies which backends support this functionality (Terraform Enterprise, Terraform Cloud, S3, GCS, Azure Storage do)

Where is State?

Local State (default)

Stored locally in a JSON format.
(`terraform.tfstate`)

Remote State

Stored on a remote source (Terraform Cloud,
Terraform Enterprise, S3, GCS, Azure Storage,
etc).

Local State

State is stored locally on one machine

It is generally acceptable for individuals and small teams

No good way to back up or share state

Tends not to scale for larger teams

Requires a more "mono repo" pattern

Remote State

State is on a remote source like S3, Terraform Cloud, Terraform Enterprise, or Consul

Remote storage is responsible for handling merging and locking

Can be queried for information in other Terraform configurations

Removes the risk of losing state to a hard drive crash or other data loss

Sensitive Data in State

State can contain sensitive data depending on the resources used

Sometimes it can contain initial database passwords or other secret data returned by a provider

Some resources support PGP encrypting the values in the state, but this is implemented on a per-resource basis

We recommend never keeping state in source control, for example

Sensitive Data in State

Local state (JSON) is not encrypted

Remote state encryption is backend-specific

State is only held in memory when remote state is used

Example: S3 bucket can be encrypted + IAM + TLS connection

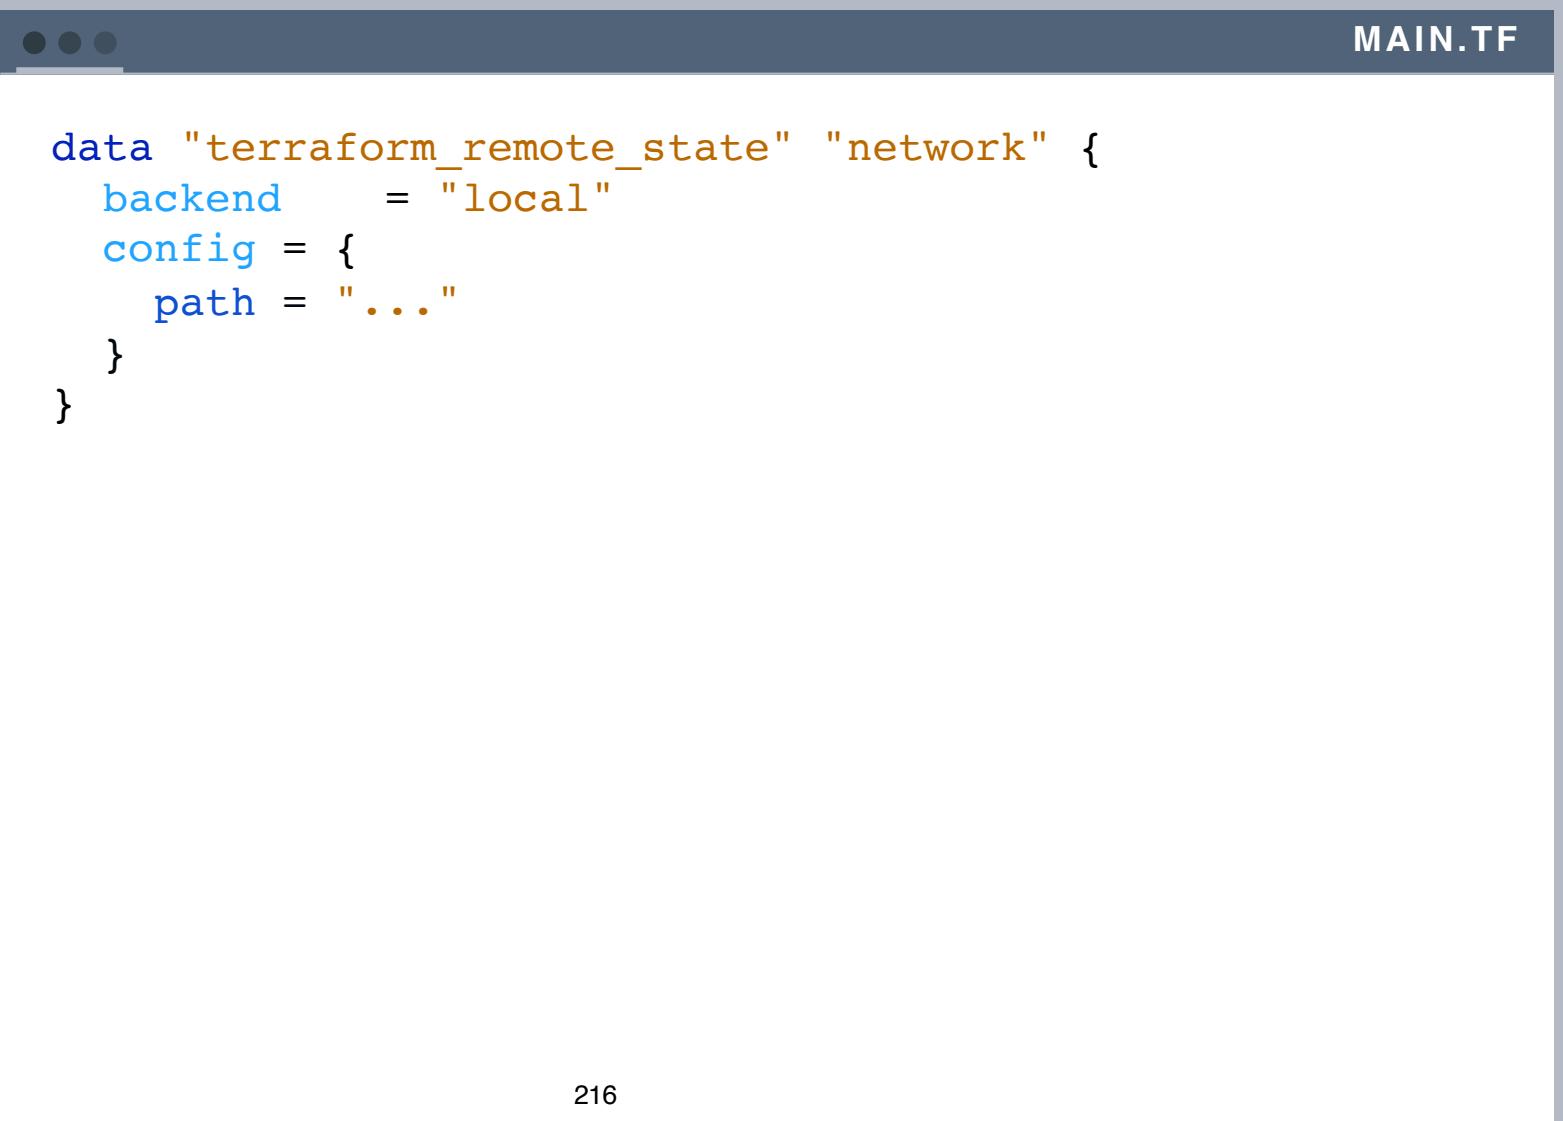
Example: TFE encrypted in transit and rest + full audit log

State as a Data Source

Terraform state data can be a data source for another configuration using `terraform_remote_state`

Only outputs are available from the remote state data

Read access to state data is required



MAIN.TF

```
data "terraform_remote_state" "network" {
    backend      = "local"
    config       = {
        path = "..."
    }
}
```

EXERCISE

Lab 3: State

Duration: 10 minutes

Goal: Read state from another Terraform project.

State Import

Terraform only manages previous or imported resources.

Terraform can import state for existing resources. If you have already spun up an instance with the web console, you may need to import it so Terraform knows about it.

This only writes the state file, not the Terraform configuration.

State Import

Critical Thinking

After importing an existing resource, what operation(s) would the terraform plan show?

State Import

Critical Thinking

After importing an existing resource, what operation(s) would the terraform plan show?

Answer

Destroy (state exists, no configuration)

EXERCISE

Lab 4: Import

Duration: 10 minutes

Goal: Use terraform import to manage existing resources.

Terraform Cloud

Terraform Ecosystem

Terraform CLI	Terraform Cloud	Terraform Enterprise
Command line tool	SaaS app	On premise/in your cloud
Infrastructure as Code	Free for individuals and teams	Workspace & team management
Open Source	Remote state	Governance and policy features
150+ Providers	Version Control System Integration	Private module registry
Module Registry		

The screenshot shows the Terraform Enterprise interface for the workspace "geoffrey-org / training-lab-dev". The "States" tab is selected. The page displays a list of recent state saves:

State ID	Saved By	Time Ago
New state #sv-Q8jFw8VkJFH6uVL7	geoffrey from Terraform	3 days ago
New state #sv-EiZEpZjAnticBGh	geoffrey from Terraform	3 days ago
New state #sv-Lzt4kV8LbNa68fXV	geoffrey from Terraform	4 days ago
New state #sv-kbJqHaEsaV6QVbgG	geoffrey from Terraform	4 days ago
New state #sv-as5VdWx1TcL7N4UH	geoffrey from Terraform	4 days ago

At the bottom of the list, there is a link: <https://atlas.hashicorp.com/app/geoffrey-org/training-lab-dev/states/sv-kbJqHaEsaV6QVbgG>.

Transitioning from Local to Remote

Transitioning is a one-time operation

After configured, Terraform will no longer store local state

The `terraform` stanza declares the configuration alongside resources



REMOTE CONFIG

```
terraform {
  cloud {
    organization = "example-company"

  # Single workspace:
    workspaces {
      name = "my-app-prod"
    }

  # Or use multiple workspaces:
  #   workspaces {
  #     tags = [ "app1" ]
  #   }
}
```



REMOTE CONFIG

```
data "terraform_remote_state" "vpc" {
    backend = "remote"

    config {
        organization = "example-org"
        hostname = "app.terraform.io"
    }

    workspaces {
        name = "example-workspace"
    }
}

resource "azurerm_network_interface" "example" {
    # ...
    subnet_id = data.terraform_remote_state.web.outputs.subnet_id
}
```



```
~/.TERRAFORMRC
```

```
credentials "app.terraform.io" {
    token = "<TOKEN>"
}
```



REMOTE CONFIG

```
data "terraform_remote_state" "vpc" {
    backend = "remote"

    config {
        organization = "example-org"
        hostname = "app.terraform.io"
        access_token = "xxxxxx"
    }
}

resource "azurerm_network_interface" "example" {
    # ...
    subnet_id = data.terraform_remote_state.web.outputs.subnet_id
}
```

Backend Initialization

Remote State backends must be initialized before use

Initialization prepares the remote backend for storage

Initialization will prompt for any missing values and cache them locally in `.terraform` (which should be ignored from source)

Initialization will also optionally migrate state

```
> terraform init
```

Initializing the backend...

Acquiring state lock. This may take a few moments...

Do you want to copy existing state to the new backend?

Pre-existing state was found while migrating the previous "local" backend to the newly configured "remote" backend. No existing state was found in the newly configured "remote" backend.

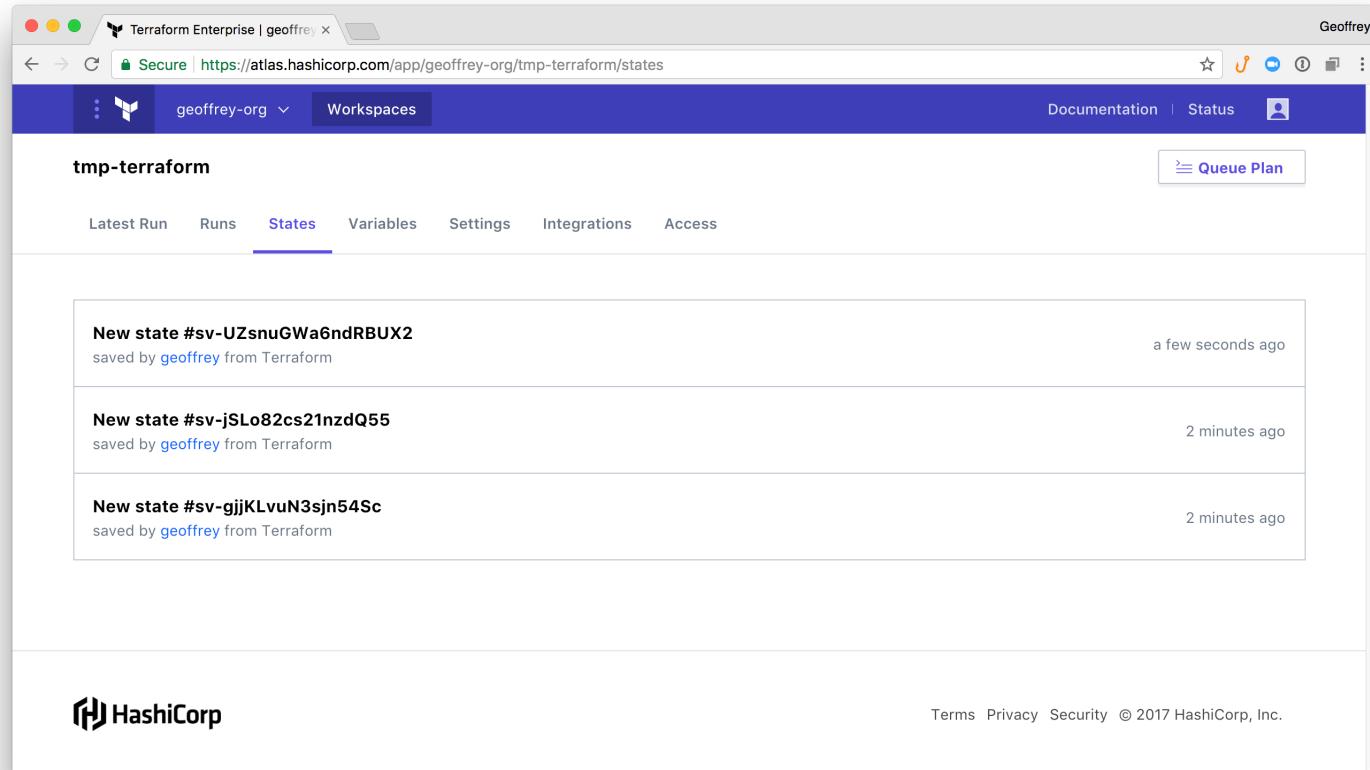
Do you want to copy this state to the new "remote" backend? Enter "yes" to copy and "no" to start with an empty state.

Enter a value:



TERMINAL

```
Successfully configured the backend "remote"! Terraform  
will automatically use this backend unless the backend  
configuration changes.
```



Remote State

Terraform will no longer write to the `.tfstate` file
Everything else about your local workflow remains
the same

TERMINAL

```
> terraform apply
azurerm_resource_group.web: Refreshing state...
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

EXERCISE

Lab 5: Terraform Cloud Remote State

Duration: 20 minutes

Goal: Use Terraform Cloud to store your state data

EXERCISE

Lab 6: Secure Variables

Duration: 20 minutes

Goal: Securely migrate and store variables to Terraform Cloud

Resource Lifecycles

What You'll Learn



- ➊ How Terraform creates and destroys resources
- ➋ How to diverge from the default behavior
- ➌ What code organization techniques go with
which lifecycle behavior

Default Behavior



MODIFY

When possible

EXAMPLES

Add a tag
Change SSH user (GCP)



DESTROY

Freely

EXAMPLES

Rename an instance
Change machine type
Decrease count
Refactor to modules
Change startup script



Do Nothing

Occasionally

EXAMPLES

Change provisioner
Add an output
Refactor to variables



terra form plan

TERMINAL

```
$ terraform plan

An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:
  ~ update in-place
  -/+ destroy and then create replacement

Terraform will perform the following actions:

  ~ aws_instance.example
    vpc_security_group_ids.#: "" => <computed>

  -/+ aws_security_group.training (new resource required)
    id: "sg-6d36fc1c" => <computed> (forces new resource)
```



Important :

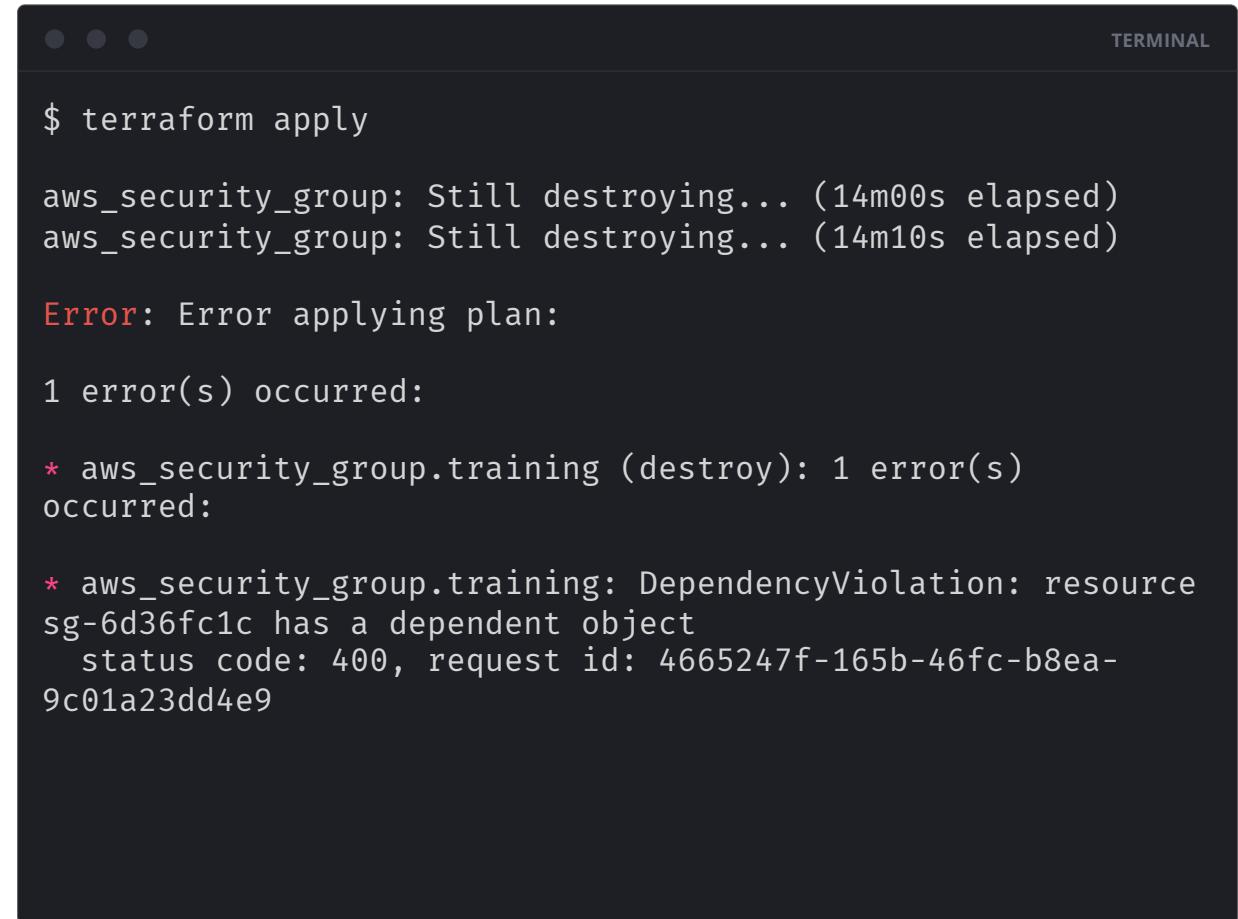
Destroy, Then Create

Existing resources are first destroyed before new ones are created.

If the destruction succeeded cleanly, then and only then are replacement resources created.



Dependency Errors



A screenshot of a terminal window titled "TERMINAL". The window shows the output of a Terraform "apply" command. It starts with three gray dots at the top, followed by the command "\$ terraform apply". Below that, two "aws_security_group" resources are shown as still destroying. An "Error" message follows, stating "Error applying plan: 1 error(s) occurred:". Two specific errors are listed, both related to an "aws_security_group.training" resource. The first error is a "DependencyViolation" due to a dependent object named "sg-6d36fc1c". The second error provides a status code of 400 and a request ID of "4665247f-165b-46fc-b8ea-9c01a23dd4e9".

```
$ terraform apply

aws_security_group: Still destroying... (14m00s elapsed)
aws_security_group: Still destroying... (14m10s elapsed)

Error: Error applying plan:

1 error(s) occurred:

* aws_security_group.training (destroy): 1 error(s) occurred:

* aws_security_group.training: DependencyViolation: resource sg-6d36fc1c has a dependent object
  status code: 400, request id: 4665247f-165b-46fc-b8ea-9c01a23dd4e9
```



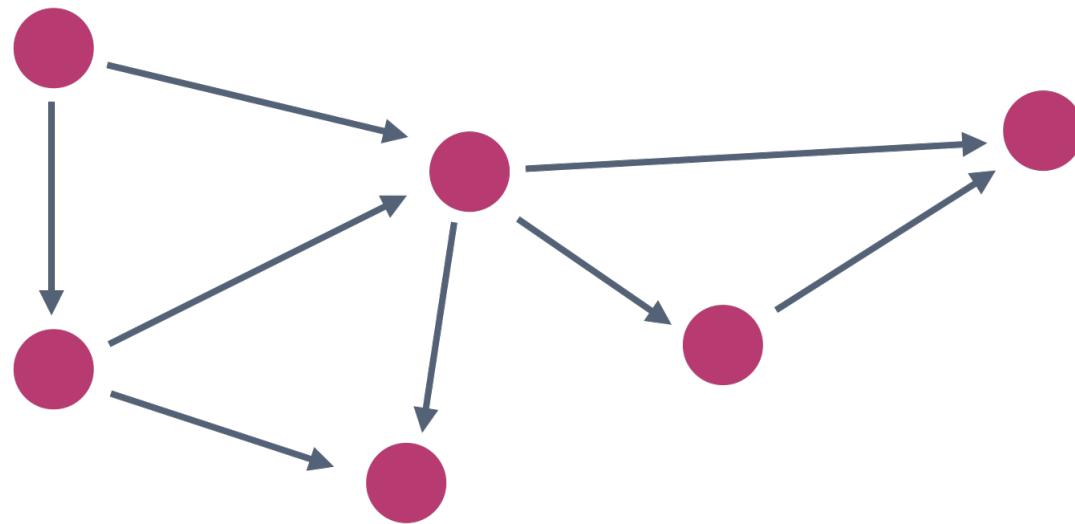
The lifecycle block

To control dependency errors

```
lifecycle {  
    create_before_destroy = true  
    prevent_destroy = true  
}
```

CODE EDITOR

Improved Application of Graph Theory



Partial Applies Will Clean Up On The Next Apply



When using `create_before_destroy`, it's possible to encounter a runtime error where Terraform cannot complete the apply.

This could leave both the old and the new instances intact.

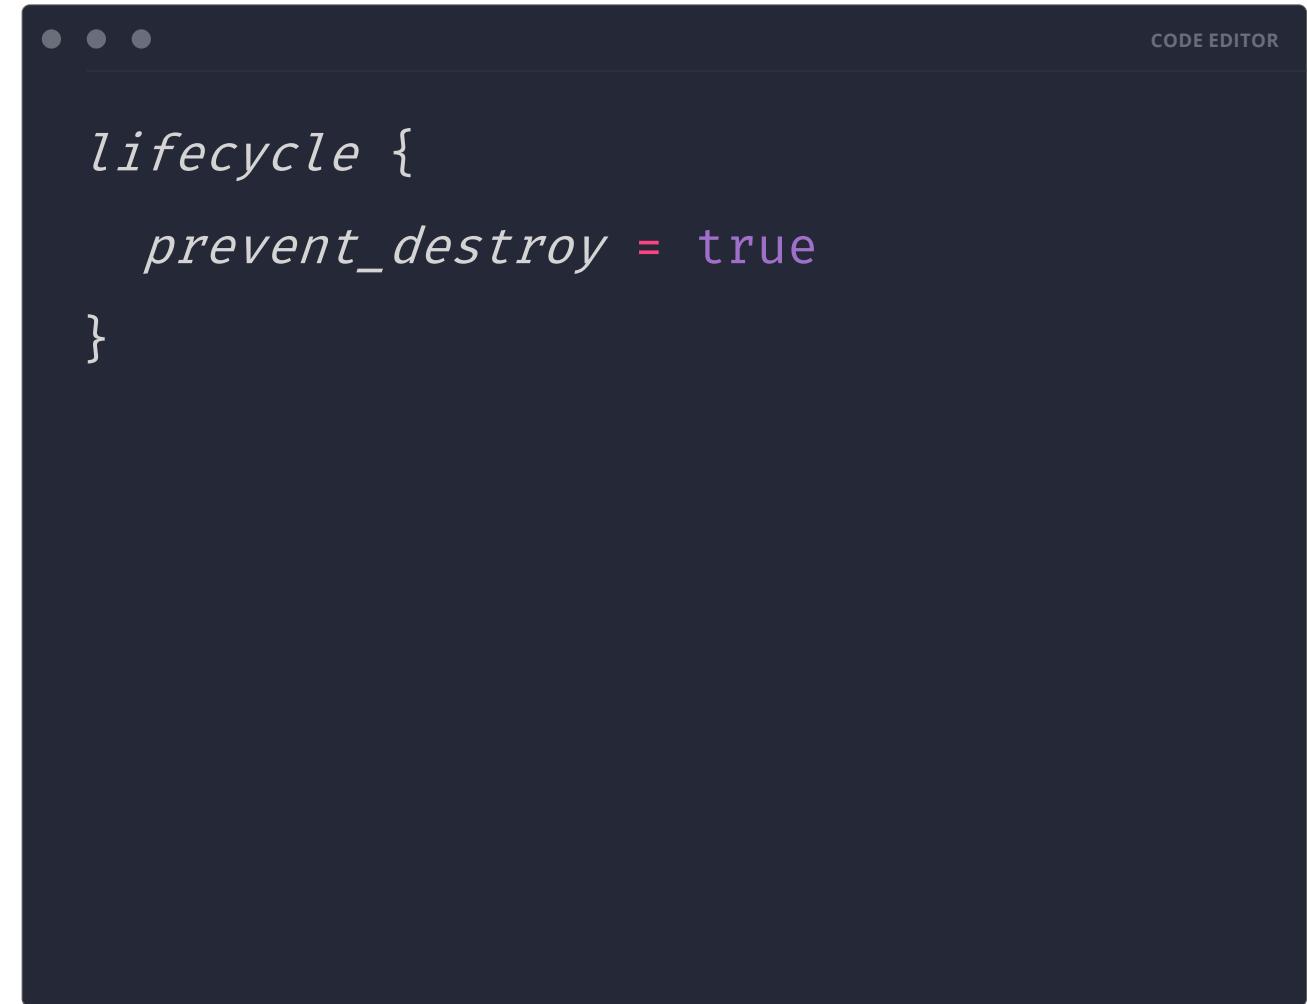
The old instance is stored in `terraform.tfstate` under a deposited key. It will be destroyed after the next successful apply.



Prevent Destroy

Warns if any change would result in destroying a resource

All resources that this resource depends on must also be set to prevent_destroy



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a single line of Terraform configuration code:

```
lifecycle {  
    prevent_destroy = true  
}
```

The code editor has three dots at the top left and a "CODE EDITOR" label at the top right. The background of the slide is white, and there is a thin horizontal line separating the slide content from the code editor window.



prevent_destroy error

Here's an example of the error message you'll see if you try to destroy a resource protected by `prevent_destroy`.

```
TERMINAL

Error: Error running plan: 1 error(s) occurred:

* google_compute_instance.test: 1 error(s)
  occurred:

* google_compute_instance.test:
  google_compute_instance.test: the plan would
  destroy this resource, but it currently has
  lifecycle.prevent_destroy set to true. To avoid
  this error and continue with the plan, either
  disable lifecycle.prevent_destroy or adjust the
  scope of the plan using the -target flag.
```

EXERCISE

Lab 7: Lifecycles

Duration: 10 minutes

Goal: Demonstrate how to apply the lifecycle directives that we've discussed so far

Template Files

What You'll Learn

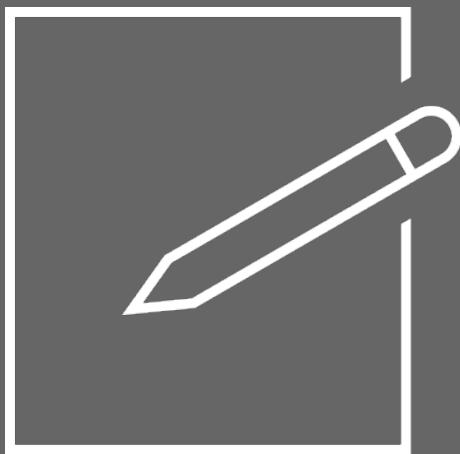


- ❖ Use the templatefile function
- ❖ Understand real world scenarios for using templates

What is a Template File in Terraform?



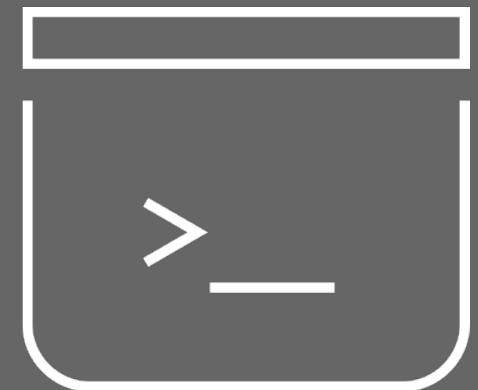
- ❖ Reads a file at a given path and renders its content as a template using a supplied set of variables
- ❖ Invoked using the keyword `templatefile`
- Uses the same interpolation syntax as strings



Plain Text
Template



Variables from Cloud or
Other Resources



Render with
Terraform

Creating a Template File



```
TERMINAL
$ mkdir -p tfproject/templates && cd $_
~/tfproject/templates
$ touch cloud-init.tpl
```

```
CODE EDITOR
#cloud-config
packages:
- azure-cli
- jq
- curl
- unzip
write_files:
- content: |
    set -e -o pipefail
    if [ -z "${TFCA_VERSION}" ]; then
```



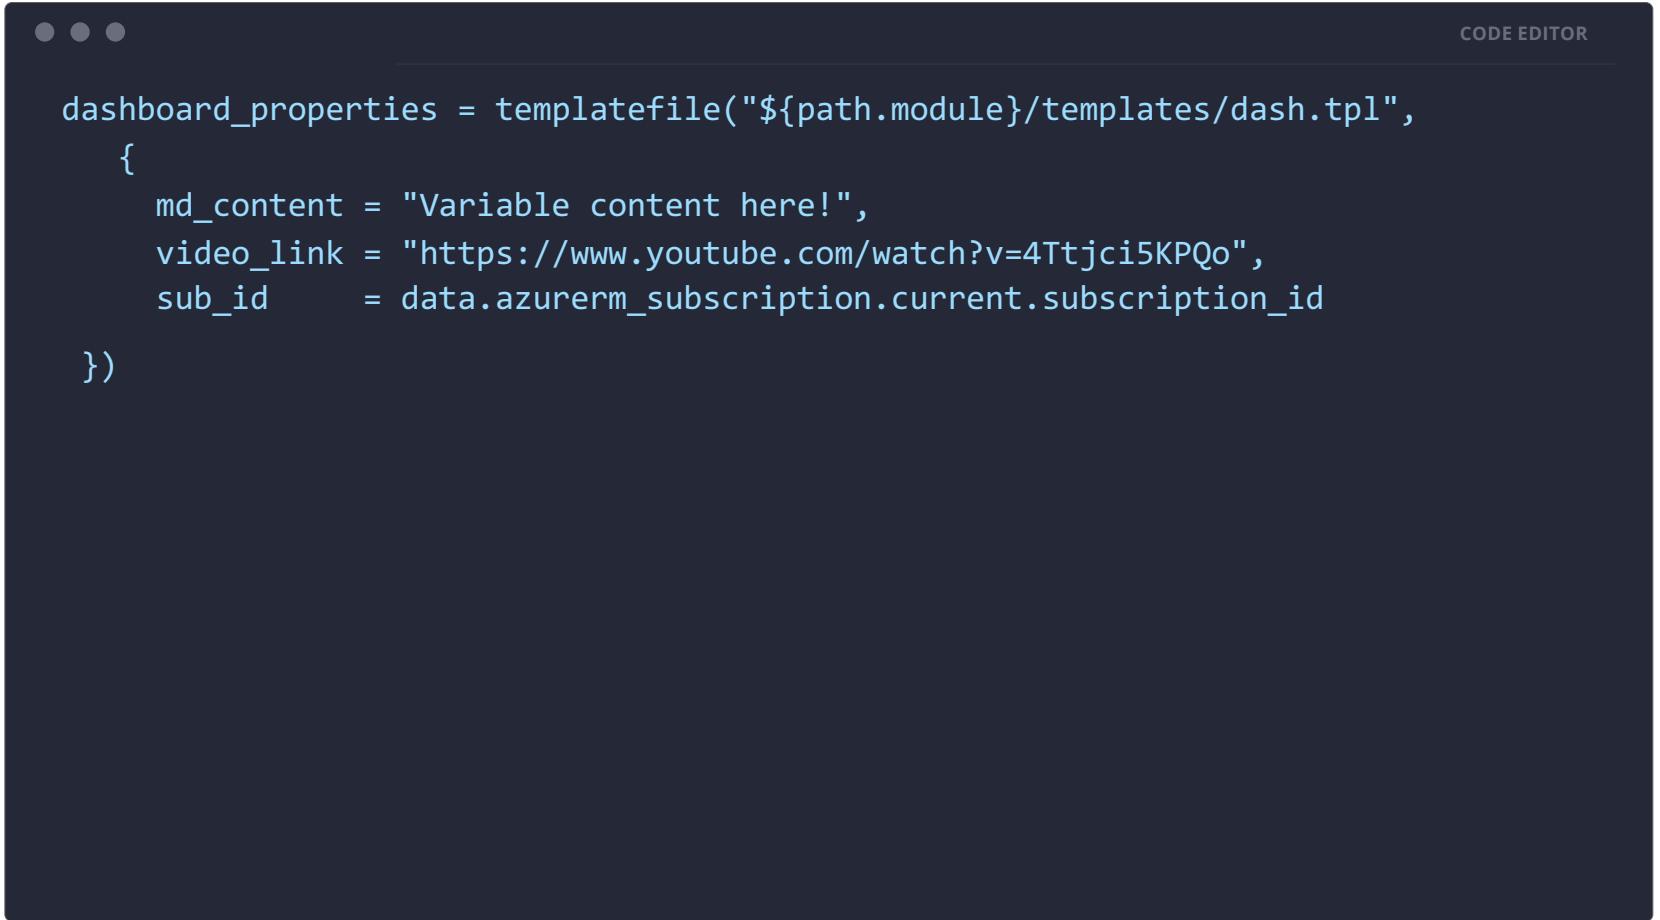
Interpolation in Template Files

We can replace values
in the template for
more dynamic
rendering

The screenshot shows a dark-themed code editor window titled "CODE EDITOR". At the top, there are three small circular icons. The main content area contains the following code:

```
• • •
● runcmd:
● - ["/bin/bash", "-c", ↓
"TFCA_VERSION='${tfca_version}'"
/opt/tfca-install/tfca-download.sh"]
```

A green downward-pointing arrow is positioned over the line "TFCA_VERSION='\${tfca_version}'" to indicate where interpolation is occurring.



A screenshot of a dark-themed code editor window. At the top right, it says "CODE EDITOR". In the top left corner, there are three small circular icons. The main area contains the following Terraform code:

```
dashboard_properties = templatefile("${path.module}/templates/dash.tpl",
{
  md_content = "Variable content here!",
  video_link = "https://www.youtube.com/watch?v=4Ttjci5KPQo",
  sub_id     = data.azurerm_subscription.current.subscription_id
})
```

Real world scenarios



Pass custom scripts to a VM's custom_data

Create a custom Azure RBAC role with

Customize a configuration file, such as a startup
script or a Consul config

Build Absolute Paths With path.module



When rendering templates from within modules, it's possible to end up in the wrong directory.

To fix this, use `path.module` when supplying the path to a template. This builds an absolute path to the current module.

```
template = templatefile("${path.module}/templates/init.tpl",
{input1 = var.input1,
input2 = var.input2 })
```



Example IAM Policy

```
resource "aws_iam_policy" "policy" {
  name = "policy"
  policy =
  templatefile("${path.module}/templates/iam_policy.tpl.json",
{
  region = var.region,
  owner_id = var.owner_id })
}
```



Example Custom Data startup script

```
resource "azurerm_linux_virtual_machine" "web" {
    name = "my-vm"

    Custom_data =
        templatefile("${path.module}/templates/user-
data.tpl.json", {username = var.username})
}
```

EXERCISE

Lab 8: Template Files

Duration: 10 minutes

Goal: Define and deploy a template file within a Terraform configuration

Debug

EXERCISE

Lab 9: Debug

Duration: 10 minutes

Goal: Set Terraform configuration to display different levels of debugging output.

Terraform Associate Certification Course

Q & A

General Discussion